



ARISTOTLE UNIVERSITY OF THESSALONIKI



FACULTY OF ENGINEERING

Pattern Recognition & Machine Learning

Bayes Decision Theory - Training

Panagiotis C. Petrantonakis

Assistant Professor

Dept. of Electrical and Computer Engineering

ppetrant@ece.auth.gr

Fall Semester

Introduction

- We saw that we can design an optimal classifier if we know the priors $P(\omega_i)$ and the class-conditional densities $p(\mathbf{x}|\omega_i)$.
- In practice, we rarely have this probabilistic structure of a problem.
- But we have...

Introduction

- We saw that we can design an optimal classifier if we know the priors $P(\omega_i)$ and the class-conditional densities $p(\mathbf{x}|\omega_i)$.
- In practice, we rarely have this probabilistic structure of a problem.
- But we have...training data!!

Introduction

- We saw that we can design an optimal classifier if we know the priors $P(\omega_i)$ and the class-conditional densities $p(x|\omega_i)$.
 - In practice, we rarely have this probabilistic structure of a problem.
 - But we have...training data!!
 - Thus, how can we use them to train our classifier?
 - Maximum likelihood estimation
 - Bayesian Estimation
 - Non-Parametric Techniques
 - First two: parameterize conditional density. Then estimate parameters.
 - Last one: end up with arbitrary distributions.
-

Independent Identically Distributed (i.i.d.) samples

- Suppose we have c datasets D_1, \dots, D_c .
- We say that: samples D_j are i.i.d. if they have been drawn independently from the same distribution $p(\mathbf{x}|\omega_j)$.
- Assuming that $p(\mathbf{x}|\omega_j)$ has a known parametric form, e.g. $p(\mathbf{x}|\omega_j) \sim \mathcal{N}(\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$:

$$p(\mathbf{x}|\omega_j) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}_j|^{1/2}} \exp \left[-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_j)^t \boldsymbol{\Sigma}_j^{-1} (\mathbf{x} - \boldsymbol{\mu}_j) \right]$$

- We can write $p(\mathbf{x}|\omega_j)$ as $p(\mathbf{x}|\omega_j, \boldsymbol{\theta}_j)$ to show the dependence of $p(\mathbf{x}|\omega_j)$ on $\boldsymbol{\theta}_j$ where $\boldsymbol{\theta}_j$ consists of the components of $\boldsymbol{\mu}_j$ and $\boldsymbol{\Sigma}_j$.

Maximum Likelihood (ML) Estimation

- **Training:** Our problem is to use the information provided by the training samples to obtain **estimates** for the unknown parameter vectors θ_j , $j = 1, \dots, c$ associated with each class.
- ML is such a method with some attractive attributes
 - good convergence: especially with large training sets
 - simple (compared to Bayesian approach and non-parametric techniques)
- To further simplify things we assume:
 - D_i gives no information about θ_j if $i \neq j$. Thus the parameters for different classes are independent.
- With the above assumption we have c separate problems of the form:
 - Use a set D of training samples, drawn i.i.d from $p(x|\theta)$, to estimate θ .

Maximum Likelihood (ML) Estimation


- Suppose D contains N samples, i.e., $\mathbf{x}_n, n = 1, \dots, N$.
- Because the samples are i.i.d. we have:

$$p(D|\boldsymbol{\theta}) = p(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N | \boldsymbol{\theta}) = \prod_{n=1}^N p(\mathbf{x}_n | \boldsymbol{\theta})$$


- Viewed as a function of $\boldsymbol{\theta}$, $p(D|\boldsymbol{\theta})$ is called the **likelihood** of $\boldsymbol{\theta}$ with respect to the samples of the dataset D (training set).
- The ML estimation of $\boldsymbol{\theta}$, i.e., $\hat{\boldsymbol{\theta}}_{ML}$, is the value of $\boldsymbol{\theta}$ that maximizes $p(D|\boldsymbol{\theta})$
- Intuitively, $\hat{\boldsymbol{\theta}}_{ML}$ best supports the observed data.

Maximum Likelihood (ML) Estimation -Example


$$p(D|\theta) = \prod_{n=1}^N p(x_n|\theta)$$



Maximum Likelihood (ML) Estimation -Example



Maximum Likelihood (ML) Estimation -Example




Duda, Hart, Stork

Maximum Likelihood (ML) Estimation -Example

- If we had a very large number of training points, what would happen to the likelihood?

$$p(D|\theta) = \prod_{n=1}^N p(x_n|\theta)$$




Duda, Hart, Stork

Maximum Likelihood (ML) Estimation -Example

- If we had a very large number of training points, what would happen to the likelihood?
 - It would be very narrow!

$$p(D|\theta) = \prod_{n=1}^N p(x_n|\theta)$$



Duda, Hart, Stork

Maximum Likelihood (ML) Estimation -Example

- If we had a very large number of training points, what would happen to the likelihood?


- It would be very **narrow!**

$$p(D|\theta) = \prod_{n=1}^N p(x_n|\theta)$$

- **Notice that:**

- the likelihood is shown as a function of θ (in this case the mean of the gaussian distribution) whereas the conditional density is shown as a function of x .

- As a likelihood, i.e., a function of θ , it is **not a probability density function** and thus its area has no significance.



Duda, Hart, Stork

Maximum Likelihood (ML) Estimation

- By definition, $\hat{\theta}_{ML}$, is the vector that maximizes likelihood:




$$p(D|\theta) = \prod_{n=1}^N p(x_n|\theta)$$

- Thus it needs to:

$$\nabla_{\theta} \prod_{n=1}^N p(x_n|\theta) = 0$$

- It is usually easier to work with logarithms.
- As the logarithm is a monotonically increasing function, if $\hat{\theta}_{ML}$ maximizes the likelihood function it also maximizes the log-likelihood, i.e., $l(\theta) = \ln p(D|\theta)$.

$$p(D|\theta) = \prod_{n=1}^N p(x_n|\theta)$$



Duda, Hart, Stork

Log-likelihood

- We have the gradient operator: $\nabla_{\theta} = \begin{bmatrix} \frac{\partial}{\partial \theta_1} \\ \frac{\partial}{\partial \theta_2} \\ \vdots \\ \frac{\partial}{\partial \theta_d} \end{bmatrix}$

- Using the log-likelihood

$$l(\theta) = \ln \prod_{n=1}^N p(x_n | \theta)$$

- We have:

$$\nabla_{\theta} l(\theta) = \nabla_{\theta} \prod_{n=1}^N p(x_n | \theta) = \sum_{n=1}^N \nabla_{\theta} \ln p(x_n | \theta) = 0$$

- Thus, we have a set of ? equations: $\nabla_{\theta} l(\theta) = \mathbf{0}$, from which we can find $\hat{\theta}_{ML}$

Log-likelihood

- We have the gradient operator: $\nabla_{\theta} = \begin{bmatrix} \frac{\partial}{\partial \theta_1} \\ \frac{\partial}{\partial \theta_2} \\ \vdots \\ \frac{\partial}{\partial \theta_d} \end{bmatrix}$

- Using the log-likelihood

$$l(\theta) = \ln \prod_{n=1}^N p(x_n | \theta)$$

- We have:

$$\nabla_{\theta} l(\theta) = \nabla_{\theta} \prod_{n=1}^N p(x_n | \theta) = \sum_{n=1}^N \nabla_{\theta} \ln p(x_n | \theta) = 0$$

- Thus, we have a set of d equations: $\nabla_{\theta} l(\theta) = \mathbf{0}$, from which we can find $\hat{\theta}_{ML}$

Maximum Likelihood (ML) Estimation

- Recap:
 - Assuming a probabilistic model for our data, e.g., $p(\mathbf{x}|\boldsymbol{\theta}) \sim \mathcal{N}(\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$, we can use the corresponding training datasets to find the optimum parameters for this model.
 - We do so by solving the system $\nabla_{\boldsymbol{\theta}} l(\boldsymbol{\theta}) = \mathbf{0}$
- Given that we have selected a **reliable model**, the ML classifier can give excellent results!
- What if the model is wrong, i.e., it does not adequately represent the underlying distributions?
 - Then the results will be probably poor!
- If the assumed model is poor we cannot assure that the classifier that we derive is the best!

Bayesian Estimation (or Bayesian Learning)

- The answers that we get from BE are nearly identical with those of ML. Nevertheless, there is a fundamental difference between the two approaches:
 - In ML the parameter's value is fixed, whereas in BE we consider θ to be a random variable.

Bayesian Estimation (or Bayesian Learning)

- The answers that we get from BE are nearly identical with those of ML. Nevertheless, there is a fundamental difference between the two approaches:
 - In ML the parameter's value is fixed, whereas in BE we consider θ to be a random variable.
- Training data allow us to convert the distribution of the parameter under consideration, i.e., $p(\theta)$, into a posterior probability.
- Bayes: $P(\omega_i|x) = \frac{p(x|\omega_i)p(\omega_i)}{p(x)}$

Bayesian Estimation (or Bayesian Learning)

- The answers that we get from BE are nearly identical with those of ML. Nevertheless, there is a **fundamental difference** between the two approaches:
 - In ML the parameter's value is fixed, whereas in BE we consider θ to be a random variable.
- Training data allow us to convert the distribution of the parameter under consideration, i.e., $p(\theta)$, into a **posterior probability**.
- Bayes: $P(\omega_i|x) = \frac{p(x|\omega_i)p(\omega_i)}{p(x)}$
- Thus, assuming that the conditional density and the a priori probabilities are unknown, how can we estimate a posteriori probabilities using training data?
- In other words, again, if we have c datasets D_1, \dots, D_c , one for each class, we want to find:

$$P(\omega_i|x, D_i) = \frac{p(x|\omega_i, D_i)p(\omega_i, D_i)}{\sum_{j=1}^c p(x|\omega_j, D_j)p(\omega_j, D_j)}$$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D_i) = \frac{p(\mathbf{x} | \omega_i, D_i) P(\omega_i | D_i)}{\sum_{j=1}^c p(\mathbf{x} | \omega_j, D_j) P(\omega_j | D_j)}$$

- We need to determine both the class-conditional densities along with the prior probabilities.
- We can assume that the priors are known or easily computed by training data. Thus the above equation becomes:

$$P(\omega_i | \mathbf{x}, D_i) = \frac{p(\mathbf{x} | \omega_i, D_i) P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x} | \omega_j, D_j) P(\omega_j)}$$

As we did in ML case, the samples of D_i have no influence on the estimation of the parameter set of $p(\mathbf{x} | \omega_j, D_j)$ if $i \neq j$.

Bayesian Estimation

- For simplicity we rewrite the equation:

$$P(\omega_i | \mathbf{x}, D_i) = \frac{p(\mathbf{x} | \omega_i, D_i) P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x} | \omega_j, D_j) P(\omega_j)}$$

as:

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x} | D) P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x} | D) P(\omega_j)}$$

where D is the training set for a particular class i and $p(\mathbf{x})$ is the probability density of that class.

- Although the desired probability density $p(\mathbf{x})$ is unknown we assume (again, as in ML) that it has a known parametric form, i.e., $p(\mathbf{x} | \boldsymbol{\theta})$.

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- Any information that we may have about $\boldsymbol{\theta}$, prior to observing our data, is contained in a known prior density $p(\boldsymbol{\theta})$.
- Once we have observed our samples (data) we can get the posterior density $p(\boldsymbol{\theta}|D)$ which we hope that is going to be sharply shaped around the true value of $\boldsymbol{\theta}$:

$$p(\boldsymbol{\theta}|D) = \frac{p(D|\boldsymbol{\theta})P(\boldsymbol{\theta})}{\int p(D|\boldsymbol{\theta})P(\boldsymbol{\theta})d\boldsymbol{\theta}}$$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- Any information that we may have about $\boldsymbol{\theta}$, prior to observing our data, is contained in a known prior density $p(\boldsymbol{\theta})$.
- Once we have observed our samples (data) we can get the posterior density $p(\boldsymbol{\theta}|D)$ which we hope that is going to be sharply shaped around the true value of $\boldsymbol{\theta}$:

$$p(\boldsymbol{\theta}|D) = \frac{p(D|\boldsymbol{\theta})P(\boldsymbol{\theta})}{\int p(D|\boldsymbol{\theta})P(\boldsymbol{\theta})d\boldsymbol{\theta}}$$

- Our goal now is to compute $p(\mathbf{x}|D)$, which is as close as we can come to knowing the real $p(\mathbf{x})$.
 - How can we accomplish that?
-

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- We take $p(\mathbf{x}|D)$ by **marginalizing** the joint density $p(\mathbf{x}, \boldsymbol{\theta}|D)$ over $\boldsymbol{\theta}$:

$$p(\mathbf{x}|D) = \int p(\mathbf{x}, \boldsymbol{\theta} | D) d\boldsymbol{\theta}$$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- We take $p(\mathbf{x}|D)$ by **marginalizing** the joint density $p(\mathbf{x}, \boldsymbol{\theta}|D)$ over $\boldsymbol{\theta}$:

$$p(\mathbf{x}|D) = \int p(\mathbf{x}, \boldsymbol{\theta} | D) d\boldsymbol{\theta}$$

- If we use the **product rule** we can rewrite the above equation as:

$$p(\mathbf{x}|D) = \int p(\mathbf{x} | \boldsymbol{\theta}, D) p(\boldsymbol{\theta}|D) d\boldsymbol{\theta}$$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- We take $p(\mathbf{x}|D)$ by **marginalizing** the joint density $p(\mathbf{x}, \boldsymbol{\theta}|D)$ over $\boldsymbol{\theta}$:

$$p(\mathbf{x}|D) = \int p(\mathbf{x}, \boldsymbol{\theta} | D) d\boldsymbol{\theta}$$

- If we use the **product rule** we can rewrite the above equation as:

$$p(\mathbf{x}|D) = \int p(\mathbf{x} | \boldsymbol{\theta}, D) p(\boldsymbol{\theta}|D) d\boldsymbol{\theta}$$

and because the selection of \mathbf{x} and D are done **independently**, we get:

$$p(\mathbf{x}|D) = \int p(\mathbf{x} | \boldsymbol{\theta}) p(\boldsymbol{\theta}|D) d\boldsymbol{\theta}$$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- We take $p(\mathbf{x}|D)$ by **marginalizing** the joint density $p(\mathbf{x}, \boldsymbol{\theta}|D)$ over $\boldsymbol{\theta}$:

$$p(\mathbf{x}|D) = \int p(\mathbf{x}, \boldsymbol{\theta} | D) d\boldsymbol{\theta}$$

- If we use the **product rule** we can rewrite the above equation as:

$$p(\mathbf{x}|D) = \int p(\mathbf{x} | \boldsymbol{\theta}, D) p(\boldsymbol{\theta}|D) d\boldsymbol{\theta}$$

and because the selection of \mathbf{x} and D are done **independently**, we get:

$$p(\mathbf{x}|D) = \int p(\mathbf{x} | \boldsymbol{\theta}) p(\boldsymbol{\theta}|D) d\boldsymbol{\theta}$$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- Thus we have the class conditional density

$$p(\mathbf{x}|D) = \int p(\mathbf{x}|\boldsymbol{\theta})p(\boldsymbol{\theta}|D)d\boldsymbol{\theta}$$

linked to the posterior density:

$$p(\boldsymbol{\theta}|D) = \frac{p(D|\boldsymbol{\theta})p(\boldsymbol{\theta})}{\int p(D|\boldsymbol{\theta})p(\boldsymbol{\theta})d\boldsymbol{\theta}}$$

where $p(D|\boldsymbol{\theta})$ is again the...

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- Thus we have the class conditional density

$$p(\mathbf{x}|D) = \int p(\mathbf{x}|\boldsymbol{\theta})p(\boldsymbol{\theta}|D)d\boldsymbol{\theta}$$

linked to the posterior density:

$$p(\boldsymbol{\theta}|D) = \frac{p(D|\boldsymbol{\theta})p(\boldsymbol{\theta})}{\int p(D|\boldsymbol{\theta})p(\boldsymbol{\theta})d\boldsymbol{\theta}}$$

where $p(D|\boldsymbol{\theta})$ is again the... likelihood $p(D|\boldsymbol{\theta}) = \prod_{n=1}^N p(\mathbf{x}_n|\boldsymbol{\theta})$

Bayesian Estimation

$$P(\omega_i | \mathbf{x}, D) = \frac{p(\mathbf{x}|D)P(\omega_i)}{\sum_{j=1}^c p(\mathbf{x}|D)P(\omega_j)}$$

- These equations constitute the formal solution to the BE problem.
- They illustrate the relation to the ML estimation approach.
 - Suppose $p(D|\theta)$ reaches a sharp peak at $\theta = \hat{\theta}$.
 - if $p(\theta) \neq 0$ at $\theta = \hat{\theta}$ and it is smooth, then $p(\theta|D)$ also peaks at that particular point
 - thus $p(\mathbf{x}|D)$ will be approximately equal to $p(\mathbf{x}|\hat{\theta})$
 - This is the result that one would take when following the ML approach.
- Thus, if $p(D|\theta)$ is **very sharp** the influence of the prior information on the uncertainty of the value of θ can be ignored.
- Nevertheless, in the general case, the Bayesian solution tells us how to use all the available information to compute $p(\mathbf{x}|D)$!

$$p(\mathbf{x}|D) = \int p(\mathbf{x}|\theta)p(\theta|D)d\theta$$

$$p(\theta|D) = \frac{p(D|\theta)p(\theta)}{\int p(D|\theta)p(\theta)d\theta}$$

$$p(D|\theta) = \prod_{n=1}^N p(x_n|\theta)$$

ML vs. BE

- Computational Complexity:
 - ML: merely differential calculus or gradient search
 - BE: complex multidimensional integration
 - Interpretability:
 - ML: solution would be easier to interpret as it returns a single model from the set a designer provided
 - BE: the solution is a weighted average of models (may end up with a model that was not initially in the designer's set, i.e., assume uniform distribution and end up with another one)
 - Performance:
 - ML: does not take into consideration information about the prior distribution of the parameters
 - BE: takes into account this information and leads to better solutions
 - In summary, there are strong theoretical arguments supporting BE, though, in practice ML is simpler and can also lead to classifiers nearly as accurate.
-

Until now...


- We talked about supervised learning of Bayes classifiers assuming that the forms of the underlying density functions are known (in some sort of parametric form)!
 - Nevertheless, the common parametric forms rarely fit real data. In particular, most pdfs are unimodal (they have single local maximum).
 - Many practical applications require multimodal density functions.
 - We will examine now *non-parametric* techniques that can lead to arbitrary shapes of distributions.
 - With non-parametric techniques we make no assumptions about the forms of the underlying distributions.
-

Density estimation

- There is a variety of density estimation methods via non-parametric approaches.
- Here we will talk about two of them:
 - Parzen windows
 - k-Nearest Neighbor
- Both approaches are variations of the *histogram* approach to density estimation
- We will start with a simple example on density approximation of a single continuous variable x using histogram approach.

Density Estimation – Histogram approach


- The x -axis (one feature) is first divided into successive bins of **length h** .
- The **probability of a sample x being located in a bin** is estimated for each of the bins.
- If N is the total number of samples and k_N of these are located inside a bin, the corresponding probability is the **frequency ratio**: $P \approx \frac{k_N}{N}$
- This approximation converges to true P as $N \rightarrow \infty$.
- The **pdf value** is assumed constant within a bin and is approximately $\hat{p}(x) = \hat{p}(\bar{x}) \approx \frac{1}{h} \frac{k_N}{N}, |x - \bar{x}| \leq \frac{h}{2}$ where \bar{x} is the midpoint of the bin.



Density Estimation – Histogram approach


$$\hat{p}(x) \approx \frac{1}{h} \frac{k_N}{N}, |x - \bar{x}| \leq \frac{h}{2}$$

- Thus, this value determines the amplitude of the histogram curve (**approximation of $p(x)$**)
- We need **relatively small h** so as the assumption that $p(x)$ is constant within the bin holds.
- $\hat{p}(x)$ converges to $p(x)$ as $N \rightarrow \infty$ and:
 - $h_N \rightarrow 0$
 - $k_N \rightarrow \infty$
 - $\frac{k_N}{N} \rightarrow 0$
- By h_N we highlight the **dependency** of h on N



Density Estimation – Histogram approach

- $\hat{p}(x)$ converges to $p(x)$ as $N \rightarrow \infty$ and:
 - $h_N \rightarrow 0$ (we already justified this one)
 - $k_N \rightarrow \infty$
 - $\frac{k_N}{N} \rightarrow 0$
- $k_N \approx PN$ and k_N tends to infinity as $N \rightarrow \infty$
- As $h_N \rightarrow 0$ also $P \rightarrow 0$, thus $\frac{k_N}{N} \rightarrow 0$
- Thus, N must be “large enough,” h_N “small enough,” and k_N “large enough” too.
- How small and how large depend on the type of the pdf function and the degree of approximation one is satisfied with.



Parzen windows

- Multidimensional case:
 - Instead of bins we have hypercubes with length of side h
- Lets define the function:

$$\varphi(\mathbf{x}_i) = \begin{cases} 1, & \forall j, |x_{ij}| \leq \frac{1}{2} \\ 0, & \text{otherwise} \end{cases}$$

where x_{ij} are the components of the feature vector $\mathbf{x}_i, i = 1, \dots, N$.

- In words: $\varphi(\mathbf{x}_i)$ is 1 only if the feature vector is within a hypercube with side $h =$

Parzen windows

- Multidimensional case:
 - Instead of bins we have hypercubes with length of side h
- Lets define the function:

$$\varphi(\mathbf{x}_i) = \begin{cases} 1, & \forall j, |x_{ij}| \leq \frac{1}{2} \\ 0, & otherwise \end{cases}$$


where x_{ij} are the components of the feature vector $\mathbf{x}_i, i = 1, \dots, N$.

- In words: $\varphi(\mathbf{x}_i)$ is 1 only if the feature vector is within a hypercube with side $h = 1$ and center at the ...

Parzen windows

- Multidimensional case:
 - Instead of bins we have hypercubes with length of side h
- Lets define the function:

$$\varphi(\mathbf{x}_i) = \begin{cases} 1, & \forall j, |x_{ij}| \leq \frac{1}{2} \\ 0, & otherwise \end{cases}$$



where x_{ij} are the components of the feature vector $\mathbf{x}_i, i = 1, \dots, N$.


- In words: $\varphi(\mathbf{x}_i)$ is 1 only if the feature vector is within a hypercube with side $h = 1$ and center at the ... origin and zero outside of the hypercube.

Parzen windows

- In the general case:

$$\varphi\left(\frac{x_i - \bar{x}}{h}\right) = \begin{cases} 1, & \forall j, |x_{ij} - \bar{x}_j| \leq \frac{h}{2} \\ 0, & \text{otherwise} \end{cases}$$

- Thus, how $\hat{p}(\bar{x}) \approx \frac{1}{h} \frac{k_N}{N}$, changes in order to account for the general case?



Theodoridis, Koutroubas

Parzen windows


- In the general case:

$$\varphi\left(\frac{x_i - \bar{x}}{h}\right) = \begin{cases} 1, & \forall j, |x_{ij} - \bar{x}_j| \leq \frac{h}{2} \\ 0, & \text{otherwise} \end{cases}$$

- Thus, how $\hat{p}(\bar{x}) \approx \frac{1}{h^d} \frac{k_N}{N}$, changes in order to account for the general case?

$$\hat{p}(x) \approx \frac{1}{h^d} \left(\frac{1}{N} \sum_{i=1}^N \varphi\left(\frac{x_i - \bar{x}}{h}\right) \right)$$

- **Interpretation:** we estimate the pdf at point \bar{x} and within a hypercube with side h . The summation equals k_N




Theodoridis, Koutroubas

Parzen windows

$$\hat{p}(x) \approx \frac{1}{h^d} \left(\frac{1}{N} \sum_{i=1}^N \varphi \left(\frac{x_i - \bar{x}}{h} \right) \right)$$

- There is an issue:
 - we try to estimate a continuous function $p(x)$ using an expansion of discontinuous terms, i.e., step function $\varphi(\cdot)$
- This led Parzen to generalize the above approximation using **smooth functions** instead of discontinuous $\varphi(\cdot)$
- Such smooth functions are known as *kernels* or **Parzen windows** and provided that $\varphi(x) \geq 0$ and $\int_x \varphi(x)dx = 1$ this approach leads to good approximations.



Theodoridis, Koutroubas

Parzen windows – gaussian kernel


- If we use the gaussian kernel the approximation for the density $p(\mathbf{x})$ would be:

$$\hat{p}(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^n \frac{1}{(2\pi)^{\frac{d}{2}} h^d} \exp\left(-\frac{(\mathbf{x} - \mathbf{x}_i)^T (\mathbf{x} - \mathbf{x}_i)}{2h^2}\right)$$


- **In words:** the unknown pdf is approximated as an *average* of N gaussians each one centered at a different point of the **training set**.
 - The **smaller the h** the narrower the kernel will be and thus each individual gaussian is more “spiky” and localized around its mean value (here the feature vector under consideration).
 - The **larger the h** the broader the kernel and the influence in the neighbor around its mean is more important.
-

Parzen windows – gaussian kernel - example


- In general we need to find the appropriate h value in accordance with the number of samples.




$h = 0.1, N = 1000$



$h = 0.1, N = 20000$



$h = 0.8, N = 1000$



$h = 0.8, N = 20000$

Theodoridis, Koutroubas

For fixed h the number of samples influence the smoothness of the resulting approximation.

For the same example, the number of samples has little influence on the approximation accuracy for a larger h

Classification using Parzen windows

- We have already seen that for the two-class classification we can minimize the overall risk by classifying ω_1 if:

$$\frac{p(\mathbf{x}|\omega_1)}{p(\mathbf{x}|\omega_2)} > \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)}$$

With the Parzen approach we can estimate the corresponding density function of the first part and the above rule becomes:

$$\frac{\frac{1}{h^d} \left(\frac{1}{N_1} \sum_{n=1}^{N_1} \varphi \left(\frac{\mathbf{x}_n - \bar{\mathbf{x}}}{h} \right) \right)}{\frac{1}{h^d} \left(\frac{1}{N_2} \sum_{n=1}^{N_2} \varphi \left(\frac{\mathbf{x}_n - \bar{\mathbf{x}}}{h} \right) \right)} > \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)}$$

- What is $\bar{\mathbf{x}}$, N_1 and N_2 in terms of the classification task?

Classification using Parzen windows

- We have already seen that for the two-class classification we can minimize the overall risk by classifying ω_1 if:

$$\frac{p(\mathbf{x}|\omega_1)}{p(\mathbf{x}|\omega_2)} > \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)}$$

With the Parzen approach we can estimate the corresponding density function of the first part and the above rule becomes:

$$\frac{\frac{1}{h^d} \left(\frac{1}{N_1} \sum_{n=1}^{N_1} \varphi \left(\frac{\mathbf{x}_n - \bar{\mathbf{x}}}{h} \right) \right)}{\frac{1}{h^d} \left(\frac{1}{N_2} \sum_{n=1}^{N_2} \varphi \left(\frac{\mathbf{x}_n - \bar{\mathbf{x}}}{h} \right) \right)} > \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)}$$

- $\bar{\mathbf{x}}$: sample to classify, N_1 : number of samples in the training set for class ω_1 and N_2 : number of samples in the training set for class ω_2

k -Nearest Neighbor estimation

- In Parzen approach the **volume** around the sample \mathbf{x} is considered **fixed**, whereas **what changes is the number of samples** that fall within this space.
- In k -NN approach the roles are reversed, i.e., we consider the number of the **samples fixed**, i.e., k , and the **volume changes** so as to determine the minimum volume needed to include k samples.
 - low density areas: big volume
 - high density areas: small volume
 - again we can think of more complex regions than just the hypercube case, e.g. hypersphere.
- For the k -NN case the estimation of the density function is known:

$$\hat{p}(\mathbf{x}) = \frac{1}{V(\mathbf{x})} \frac{k}{N}$$

Classification using k-NN estimation

- During a classification of a test sample x we:
 - compute its distance (e.g., Euclidian) from all the training vectors (all classes, e.g., ω_1, ω_2).
 - determine r_1 : radius of the hypersphere, centered at x , that contains k samples of class ω_1
 - determine r_2 : radius of the hypersphere, centered at x , that contains k samples of class ω_2
 - estimate respective volumes, V_1 and V_2
 - assign x to class ω_1 if:

$$\frac{\frac{1}{V_1} \frac{k}{N_1}}{\frac{1}{V_2} \frac{k}{N_2}} > \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)} \rightarrow \frac{V_2}{V_1} > \frac{N_1}{N_2} \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)}$$

- If we use Mahalanobis distance we should replace hyperspheres with...

Classification using k-NN estimation


- During a classification of a test sample x we:
 - compute its distance (e.g., Euclidian) from all the training vectors (all classes, e.g., ω_1, ω_2).
 - determine r_1 : radius of the hypersphere, centered at x , that contains k samples of class ω_1
 - determine r_2 : radius of the hypersphere, centered at x , that contains k samples of class ω_2
 - estimate respective volumes, V_1 and V_2
 - assign x to class ω_1 if:

$$\frac{\frac{1}{V_1} \frac{k}{N_1}}{\frac{1}{V_2} \frac{k}{N_2}} > \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)} \rightarrow \frac{V_2}{V_1} > \frac{N_1}{N_2} \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)}$$

- If we use Mahalanobis distance we should replace hyperspheres with...hyperellipsoids.

Classification using 5-NN estimation - example

- In which class would you assign the sample denoted by a star in the following figure with classes ω_1 (black dots) and ω_2 (blue dots) being equiprobable (assume simple costs)?




$$\frac{V_2}{V_1} > \frac{N_1}{N_2} \frac{\lambda_{12} - \lambda_{22}}{\lambda_{21} - \lambda_{11}} \frac{P(\omega_2)}{P(\omega_1)}$$

Theodoridis, Koutroubas

k -Nearest Neighbor Classifier

- A variation of the previous approach of k-NN based density estimation can lead to a suboptimal but useful, in practice, **non-linear classification algorithm**: the k -NN classifier
- The **k -NN algorithm** is summarized as follows:
 - From the N training samples identify k nearest neighbors (irrespectively of the class they belong to)
 - From these k samples, find the number of samples k_i that belong to class $\omega_i, i = 1, \dots, c$.
 - Assign x to the class with the maximum k_i .
- Various distance measures can be used
- Simplest form: $k = 1$, i.e., 1-NN classifier.



Theodoridis, Koutroubas

Error using k-NN classifier

- It can be shown that as $N \rightarrow \infty$ the classification error probability for the 1-NN classifier is bounded:

$$P_B \leq P_{1-NN} \leq \dots \leq 2P_B$$

- Thus, the error committed by the 1-NN classifier is at most twice the optimal Bayes classifier.
- Moreover, the asymptotic performance of the k-NN classifier is better than 1-NN. For the two class case it can be shown that:

$$P_B \leq P_{k-NN} \leq P_B + \sqrt{\frac{2P_{1-NN}}{k}}$$

suggesting that for $k \rightarrow \infty$ the k-NN classifier tends to be optimal.

- Finally, for small Bayesian error it can be proven that: $P_{3-NN} \approx P_B + 3P_B^2$
 - Thus for $P_B = 0.01$ I get $P_{3-NN} \approx 0.0103$

Naïve Bayes Classifier

- In this lecture we presented various techniques for the estimation of the probability density function $p(x|\omega_i), i = 1, \dots, c$ based on the available training set.
 - It may be obvious until now that the number of training samples plays crucial role for the abovementioned estimation.
 - Due to *curse of dimensionality*, the bigger the number of dimensions the bigger the number of training samples needed for effective estimation.
 - Crudely speaking, if we need N samples for good estimation in one dimension, in d dimensions I would need N^d .
 - Thus, big d makes good estimation of $p(x|\omega_i)$ unrealistic.
 - We could make some assumptions to alleviate this problem.
-

Naïve Bayes Classifier

- One widely used approach is to assume that **individual features are independent**.
- Under this assumption we can write:

$$p(\mathbf{x}|\omega_i) = \prod_{j=1}^d p(x_j|\omega_i), i = 1, \dots, c$$

- Now, we only need to estimate d **one-dimensional** densities for each one of the c classes.
- This leads to the so called **Naïve Bayes** classifier which can be **very robust** to the violation of the *independence* assumption and deliver good results in many applications in practice.

Sum up...

- We saw various training techniques for the Bayes classifier
 - In essence, training is just the **estimation of the probability density** that generate my data using the available training data
 - Training approaches
 - Maximum likelihood
 - Bayesian Estimation (Bayesian Learning)
 - Non-Parametric Techniques
 - Parzen windows
 - k-NN estimation
 - We also described an alteration of k -NN estimation that leads to the celebrated k -NN classifier
 - Finally, we presented the Naïve Bayes approach that assumes different features are statistically independent.
-

Questions!

- Which aspect of the Bayes classifier is responsible for its inductive bias?

Questions!

- Which aspect of the Bayes classifier is responsible for its inductive bias?
 - probability density functions!
 - You, as a designer, can choose, these pdfs and as a result develop classifiers that have inductive biases based on what you believe/know about the problem under consideration!

Questions!

- Which aspect of the Bayes classifier is responsible for its inductive bias?
 - probability density functions!
 - You, as a designer, can choose, these pdfs and as a result develop classifiers that have inductive biases based on what you believe/know about the problem under consideration!
- Which aspect of the k -NN classifier is responsible for its inductive bias?

Questions!

- Which aspect of the Bayes classifier is responsible for its inductive bias?
 - probability density functions!
 - You, as a designer, can choose, these pdfs and as a result develop classifiers that have inductive biases based on what you believe/know about the problem under consideration!
- Which aspect of the k -NN classifier is responsible for its inductive bias?
 - label for an example should be similar to the label of nearby points.
 - all features are equally important (remember the difference with the DT)
 - if you have data with only a few relevant features and lots of irrelevant features, k -NN is likely to do poorly!



ARISTOTLE UNIVERSITY OF THESSALONIKI



FACULTY OF ENGINEERING

Questions?

Pattern Recognition & Machine Learning

Bayes Decision Theory - Training