



ARISTOTLE UNIVERSITY OF THESSALONIKI



FACULTY OF ENGINEERING

Pattern Recognition & Machine Learning

Linear Discriminant Functions and Models

Panagiotis C. Petrantonakis

Assistant Professor

Dept. of Electrical and Computer Engineering

ppetrant@ece.auth.gr

Fall Semester

Until now...

- We have seen only **one** out of **three** different approaches of solving decision problems:
 - *Generative probabilistic models*:
 - solve the inference problem of determining the class-conditional density function $p(x|\omega_i)$ and the priors $P(\omega_i)$ and then use the Bayes formula to find posterior class probabilities.
 - We used training data to infer $p(x|\omega_i)$.

$$P(\omega_i | x) = \frac{p(x|\omega_i)P(\omega_i)}{p(x)}$$

Until now...

- We have seen only **one** out of **three** different approaches of solving decision problems:
 - *Generative probabilistic models*:
 - solve the inference problem of determining the class-conditional density function $p(x|\omega_i)$ and the priors $P(\omega_i)$ and then use the Bayes formula to find posterior class probabilities.
- We used training data to infer $p(x|\omega_i)$.
- We have also talked about *discriminant functions*
 - In the above case, we replaced $P(\omega_i | x)$ with $y_i(x)$ (in the BDT lecture we use the symbol g instead of y) and under certain assumptions about the class-conditional density $p(x|\omega_i)$ we ended up with linear or non-linear *decision boundaries*.
 - We say that the corresponding discriminant functions are *linear or non-linear discriminants!*

Until now...

- We have seen only **one** out of **three** different approaches of solving decision problems:
 - *Generative probabilistic models*:
 - solve the inference problem of determining the class-conditional density function $p(\mathbf{x}|\omega_i)$ and the priors $P(\omega_i)$ and then use the Bayes formula to find posterior class probabilities.
- We used training data to infer $p(\mathbf{x}|\omega_i)$.
- We have also talked about *discriminant functions*
 - In the above case, we replaced $P(\omega_i | \mathbf{x})$ with $y_i(\mathbf{x})$ (in the BDT lecture we use the symbol g instead of y) and under certain assumptions about the class-conditional density $p(\mathbf{x}|\omega_i)$ we ended up with linear or non-linear *decision boundaries*.
 - We say that the corresponding discriminant functions are *linear or non-linear discriminants!*
- In this lecture we will talk solely about linear discriminants of the form

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

Today...

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

- We will see the rest **two** approaches of solving decision problems
 - *Linear discriminant functions* that map each input \mathbf{x} directly onto a class label.
 - For instance, in two-class classification problems $y(\mathbf{x})$ takes binary values such that $y(\mathbf{x}) = 0$ represents class ω_1 and $y(\mathbf{x}) = 1$ represents class ω_2 (probabilistic modeling plays no role in this scheme)

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

Today...

- We will see the rest **two** approaches of solving decision problems
 - *Linear discriminant functions* that map each input \mathbf{x} directly onto a class label.
 - For instance, in two-class classification problems $y(\mathbf{x})$ takes binary values such that $y(\mathbf{x}) = 0$ represents class ω_1 and $y(\mathbf{x}) = 1$ represents class ω_2 (probabilistic modeling plays no role in this scheme)
 - *Discriminative probabilistic models*
 - solve the inference problem of determining the posterior class probabilities $P(\omega_i | \mathbf{x})$ and then assign each sample \mathbf{x} to one of the classes

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

Today...

- We will see the rest **two** approaches of solving decision problems
 - *Linear discriminant functions* that map each input \mathbf{x} directly onto a class label.
 - For instance, in two-class classification problems $y(\mathbf{x})$ takes binary values such that $y(\mathbf{x}) = 0$ represents class ω_1 and $y(\mathbf{x}) = 1$ represents class ω_2 (probabilistic modeling plays no role in this scheme)
 - *Discriminative probabilistic models*
 - solve the inference problem of determining the posterior class probabilities $P(\omega_i|\mathbf{x})$ and then assign each sample \mathbf{x} to one of the classes
 - For the former approach we will see the algorithms:
 - Least squares
 - Perceptron
 - For the latter approach we will see the algorithm:
 - Logistic Regression

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$


Today...

- We will see the rest **two** approaches of solving decision problems
 - *Linear discriminant functions* that map each input \mathbf{x} directly onto a class label.
 - For instance, in two-class classification problems $y(\mathbf{x})$ takes binary values such that $y(\mathbf{x}) = 0$ represents class ω_1 and $y(\mathbf{x}) = 1$ represents class ω_2 (probabilistic modeling plays no role in this scheme)
 - *Discriminative probabilistic models*
 - solve the inference problem of determining the posterior class probabilities $P(\omega_i|\mathbf{x})$ and then assign each sample \mathbf{x} to one of the classes
 - For the former approach we will see the algorithms:
 - Least squares
 - Perceptron
 - For the latter approach we will see the algorithm:
 - Logistic Regression
 - In **all algorithms** that we will see in this lecture, we use training data to determine vector \mathbf{w} and w_0 . (**big difference with the probabilistic generative models!**)

Linear planes – Two classes

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$


- We assign every input vector to ω_1 if $y(\mathbf{x}) \geq 0$ and to ω_2 otherwise.



Linear planes – Two classes

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$


- We assign every input vector to ω_1 if $y(\mathbf{x}) \geq 0$ and to ω_2 otherwise.
- The decision boundary (or surface) is defined by $y(\mathbf{x}) = 0$.
 - for a d -dimensional feature space I get a $(d - 1)$ -dimensional (hyper) plane.



Linear planes – Two classes

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

- We assign every input vector to ω_1 if $y(\mathbf{x}) \geq 0$ and to ω_2 otherwise.
- The decision boundary (or surface) is defined by $y(\mathbf{x}) = 0$.
 - for a d -dimensional feature space I get a $(d - 1)$ -dimensional (hyper) plane.
- \mathbf{w} is the weight vector and w_0 is the bias (or threshold)
- Can you tell how \mathbf{w} is related to every vector on the decision boundary $y(\mathbf{x}) = 0$?



Linear planes – Two classes


$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

- If we consider two vectors \mathbf{x}_a and \mathbf{x}_b that both lie on the decision boundary then:

$$\mathbf{w}^t \mathbf{x}_a + w_0 = \mathbf{w}^t \mathbf{x}_b + w_0 \Leftrightarrow$$

$$\mathbf{w}^t (\mathbf{x}_a - \mathbf{x}_b) = 0$$

- Thus, \mathbf{w} is orthogonal to the decision surface and to every vector on it.
 - Determines the orientation of the decision surface!
- What is the interpretation of the bias term w_0 ?



Linear planes – Two classes


$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

- If we consider a vector \mathbf{x} that lie on the decision boundary then, $y(\mathbf{x}) = 0$ and:

$$\mathbf{w}^t \mathbf{x} + w_0 = 0 \Leftrightarrow$$

$$\frac{\mathbf{w}^t \mathbf{x}}{\|\mathbf{w}\|} = \frac{-w_0}{\|\mathbf{w}\|}$$


- Thus the bias term is actually the (signed) distance of the decision surface from the origin.
- It shifts the boundary away from the origin.



Linear planes – Two classes

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$

- Finally what is the interpretation of the $y(\mathbf{x})$ value?



Linear planes – Two classes

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$


- Finally what is the interpretation of the $y(\mathbf{x})$ value?
- If \mathbf{x} is an arbitrary sample vector and \mathbf{x}_k is its projection on the decision surface then:

$$\mathbf{x} = \mathbf{x}_k + r \frac{\mathbf{w}}{\|\mathbf{w}\|} \Leftrightarrow$$

$$\mathbf{w}^T \mathbf{x} + w_0 = \mathbf{w}^T \mathbf{x}_k + r \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + w_0 \Leftrightarrow$$

$$y(\mathbf{x}) = r \frac{\|\mathbf{w}\|^2}{\|\mathbf{w}\|} \Leftrightarrow r = \frac{y(\mathbf{x})}{\|\mathbf{w}\|}$$

- Thus, $y(\mathbf{x})$ determines the distance of the sample from the decision surface.



Least Squares

- We will see now how can we use the Least Squares approach to estimate w and w_0 (training!).

Least Squares

- We will see now how can we use the Least Squares approach to estimate \mathbf{w} and w_0 (training!).
- Assume we have c classes and N training samples. Each one of the samples is associated with a target (label) which is structured using one-hot encoding.
 - For instance, if we have 5 classes and a sample x_n has $label = 3$ then the target vector using one-hot encoding is $t_n = [0 \ 0 \ 1 \ 0 \ 0]$. We can gather all such target encoding in one matrix \mathbf{T} , where every row is the one-hot encoding of the target of the respective sample, i.e. $\mathbf{T} \in \mathbb{R}^{N \times c}$.

Least Squares

- We will see now how can we use the Least Squares approach to estimate \mathbf{w} and w_0 (training!).
- Assume we have c classes and N training samples. Each one of the samples is associated with a target (label) which is structured using one-hot encoding.
 - For instance, if we have 5 classes and a sample \mathbf{x}_n has *label* = 3 then the target vector using one-hot encoding is $\mathbf{t}_n = [0 \ 0 \ 1 \ 0 \ 0]$. We can gather all such target encoding in one matrix \mathbf{T} , where every row is the one-hot encoding of the target of the respective sample, i.e. $\mathbf{T} \in \mathbb{R}^{N \times c}$.
- Each of the c classes has its own linear discriminant:

$$y_i(\mathbf{x}) = \mathbf{w}_i^t \mathbf{x} + w_{i0}$$

- For simplicity we will adopt the shorter notation:

$$y_i(\mathbf{x}) = \tilde{\mathbf{w}}_i^t \tilde{\mathbf{x}}, \text{ where } \tilde{\mathbf{w}}_i = \begin{bmatrix} w_{i0} \\ w_{i1} \\ \vdots \\ w_{id} \end{bmatrix} \text{ and } \tilde{\mathbf{x}} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}$$

Least Squares

$$y_i(\boldsymbol{x}) = \boldsymbol{w}_i^t \boldsymbol{x}$$

Note: for the rest of the lecture we will use $\boldsymbol{w}, \boldsymbol{x}$ but we will mean $\tilde{\boldsymbol{w}}$ and $\tilde{\boldsymbol{x}}$!

Least Squares

$$y_i(\mathbf{x}) = \mathbf{w}_i^t \mathbf{x}$$

Note: for the rest of the lecture we will use \mathbf{w}, \mathbf{x} but we will mean $\tilde{\mathbf{w}}$ and $\tilde{\mathbf{x}}$!

- We can use matrix notation to include all the discriminants:

$$\mathbf{y}(\mathbf{x}) = \mathbf{W}^T \mathbf{x}$$

where the columns of \mathbf{W} contain the vectors \mathbf{w}_i for each class, i.e. $\mathbf{W} \in \mathbb{R}^{(d+1) \times c}$, and

$$\mathbf{y}(\mathbf{x}) = \begin{bmatrix} y_1(\mathbf{x}) \\ y_2(\mathbf{x}) \\ \vdots \\ y_c(\mathbf{x}) \end{bmatrix}$$

- Thus the classification problem is formed as:

$$\text{assign } \mathbf{x} \text{ to class } \omega_i \text{ if } i = \operatorname{argmax}_j y_j(\mathbf{x})$$

Least Squares

$$\mathbf{y}(\mathbf{x}) = \mathbf{W}^T \mathbf{x}$$

- Now that we have formulated our problem we only need to determine matrix \mathbf{W} .

Least Squares

$$\mathbf{y}(\mathbf{x}) = \mathbf{W}^T \mathbf{x}$$

- Now that we have formulated our problem we only need to determine matrix \mathbf{W} .
- Our objective is to make, for each sample \mathbf{x}_n , as minimum as possible the difference:

$$\mathbf{W}^T \mathbf{x}_n - \mathbf{t}_n^T$$

Least Squares

$$\mathbf{y}(\mathbf{x}) = \mathbf{W}^T \mathbf{x}$$

- Now that we have formulated our problem we only need to determine matrix \mathbf{W} .
- Our objective is to make, for each sample \mathbf{x}_n , as minimum as possible the difference:

$$\mathbf{W}^T \mathbf{x}_n - \mathbf{t}_n^T$$

if we use the matrices \mathbf{T} and \mathbf{X} , where every row of \mathbf{X} is a sample of the training set, i.e. $\mathbf{X} \in \mathbb{R}^{N \times (d+1)}$, then our optimization problem can be described as the minimization of the sum-of-squares error function:

$$E(\mathbf{W}) = \frac{1}{2} \text{tr}[(\mathbf{X}\mathbf{W} - \mathbf{T})^T (\mathbf{X}\mathbf{W} - \mathbf{T})]$$

Least Squares

- In order to minimize the sum-of-squares error function (least squares approach):

$$E(\mathbf{W}) = \frac{1}{2} \text{tr}[(\mathbf{X}\mathbf{W} - \mathbf{T})^T(\mathbf{X}\mathbf{W} - \mathbf{T})]$$

We will find the respective derivative and set equal to zero:




$$\frac{\partial E(\mathbf{W})}{\partial \mathbf{W}} = 0$$

we will come up with the solution:

$$\mathbf{W}_{LS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{T} = \mathbf{X}^\dagger \mathbf{T}$$

- Classification using the discriminant function: $\mathbf{y}_{LS}(\mathbf{x}) = \mathbf{W}_{LS} \mathbf{x}$

Least Squares – Problems!



Bishop

- The decision boundaries are very sensitive to outliers.
- for $c > 2$ there are masking phenomena (decision regions become very small or completely ignored).
- $y_{LS}(x)$ are not real probabilities and sometimes can have values outside of the range [0,1].
- Very costly to compute!

Perceptron

- Perceptron algorithm is **only for two-class** problems.

Perceptron

- Perceptron algorithm is only for two-class problems.
- Assuming input $x \in \mathbb{R}^{d+1}$ (again we consider the augmented samples $[1, x]$)

Perceptron

- Perceptron algorithm is only for two-class problems.
- Assuming input $x \in \mathbb{R}^{d+1}$ (again we consider the augmented samples $[1, x]$)
- We also assume that we have targets $t \in \{\omega_1, \omega_2\} \rightarrow t \in \{-1, 1\}$

Perceptron

- Perceptron algorithm is only for two-class problems.
- Assuming input $\mathbf{x} \in \mathbb{R}^{d+1}$ (again we consider the augmented samples $[1, \mathbf{x}]$)
- We also assume that we have targets $t \in \{\omega_1, \omega_2\} \rightarrow t \in \{1, -1\}$
- Thus, we can make a prediction based on outcome: $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$ where
$$f(a) = \begin{cases} 1, & a > 0 \\ -1, & a \leq 0 \end{cases}$$
- With this adoption I decide to assign to class ω_1 if $\mathbf{w}^T \mathbf{x} > 0$ and to class ω_2 otherwise.

$$t \in \{\omega_1, \omega_2\} \rightarrow t \in \{-1, 1\}$$

$$\omega_1 \text{ if } \mathbf{w}^T \mathbf{x} \geq 0$$

Perceptron - Training

- Assume we have N training samples, $x_n, n = 1, \dots, N$ with the respective targets (labels) t_n .
- Objective: find \mathbf{w} such that for all (x_n, t_n) pairs I get:

$$\mathbf{w}^T \mathbf{x}_n t_n > 0$$

- What is the objective function that we should minimize?

$$t \in \{\omega_1, \omega_2\} \rightarrow t \in \{-1, 1\}$$

$$\omega_1 \text{ if } \mathbf{w}^T \mathbf{x} \geq 0$$

Perceptron - Training

- Assume we have N training samples, $x_n, n = 1, \dots, N$ with the respective targets (labels) t_n .
- Objective: find \mathbf{w} such that for all (x_n, t_n) pairs I get:

$$\mathbf{w}^T \mathbf{x}_n t_n > 0$$

- What is the objective function that we should minimize?

$$E_P(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^T \mathbf{x}_n t_n$$

where $\mathcal{M}: \{n: \mathbf{w}^T \mathbf{x}_n t_n \leq 0\}$.

Perceptron - Training

$$E_P(\mathbf{w}) = -\sum_{n \in \mathcal{M}} \mathbf{w}^t \mathbf{x}_n t_n \text{ where } \mathcal{M}: \{n: \mathbf{w}^t \mathbf{x}_n t_n \leq 0\}$$

- $E_P(\mathbf{w})$ is piecewise linear. Why?

Perceptron - Training

$$E_P(\mathbf{w}) = -\sum_{n \in \mathcal{M}} \mathbf{w}^t \mathbf{x}_n t_n \text{ where } \mathcal{M}: \{n: \mathbf{w}^t \mathbf{x}_n t_n \leq 0\}$$

- $E_P(\mathbf{w})$ is piecewise linear. Why?
 - If we change \mathbf{w} linearly then $E_P(\mathbf{w})$ changes also linearly until the point where there is a change in the number of misclassified samples!
 - At these points the partial derivative of $E_P(\mathbf{w})$ cannot be defined and, thus, the derivative of $E_P(\mathbf{w})$ is a discontinuous function!
- Thus, how shall we proceed with the training of the perceptron algorithm?

Perceptron - Training

$$E_P(\mathbf{w}) = -\sum_{n \in \mathcal{M}} \mathbf{w}^t \mathbf{x}_n t_n \text{ where } \mathcal{M}: \{n: \mathbf{w}^t \mathbf{x}_n t_n \leq 0\}$$

- $E_P(\mathbf{w})$ is piecewise linear. Why?
 - If we change \mathbf{w} linearly then $E_P(\mathbf{w})$ changes also linearly until the point where there is a change in the number of misclassified samples!
 - At these points the partial derivative of $E_P(\mathbf{w})$ cannot be defined and, thus, the derivative of $E_P(\mathbf{w})$ is a discontinuous function!
- Thus, how shall we proceed with the training of the perceptron algorithm?
- Gradient Descent!

Gradient Descent

- The gradient of a function encodes all directional derivatives.
 - It points towards the direction of the steepest ascent.
 - the magnitude of the gradient is the rate of increase in that direction.
 - It is always perpendicular the contours of a function.

Gradient Descent

- The gradient of a function encodes all directional derivatives.
 - It points towards the direction of the steepest ascent.
 - the magnitude of the gradient is the rate of increase in that direction.
 - It is always perpendicular the contours of a function.
- Whenever the loss is a sum of error terms for each datapoint we can use Gradient Descent, e.g.:

$$E(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^t \mathbf{x}_n t_n = - \sum_{n \in \mathcal{M}} E_n(\mathbf{w})$$

- In general we have: $E(\mathbf{w}) = \sum_{n=1}^N E_n(\mathbf{w})$

Stochastic gradient descent

- **Gradient descent:** $\mathbf{w}^{i+1} = \mathbf{w}^i - \eta \nabla E(\mathbf{w}^i)$
 - where $E(\mathbf{w}^i) = \sum_{n=1}^N E_n(\mathbf{w}^i)$
- We can extend the above concept by adding stochasticity in our Gradient Descent algorithm (Stochastic Gradient Descent):

$$\mathbf{w}^{i+1} = \mathbf{w}^i - \eta \nabla E_n(\mathbf{w}^i)$$

or

$$\mathbf{w}^{i+1} = \mathbf{w}^i - \eta \nabla \sum_{l=1}^L E_l(\mathbf{w}^i) \text{ (*minibatch*)}$$

- Either with single points or minibatches what I compute is an approximation of the full gradient.

Stochastic Gradient Descent - Algorithm

- SGD steps:
 - Initialize w^0 , choose η
 - Iterate over datapoints: $w^{i+1} = w^i - \eta \nabla \sum_{l=1}^L E_l(w^i)$

Stochastic Gradient Descent - Algorithm

- SGD steps:
 - Initialize w^0 , choose η
 - Iterate over datapoints: $w^{i+1} = w^i - \eta \nabla \sum_{l=1}^L E_l(w^i)$
- If E is convex in w and η is small enough SGD will converge
 - very small η slow convergence
 - large η : probably it will not converge
- Scheduling for η is important (maybe decrease η with iterations)


Stochastic Gradient Descent - Algorithm

- SGD steps:
 - Initialize w^0 , choose η
 - Iterate over datapoints: $w^{i+1} = w^i - \eta \nabla \sum_{l=1}^L E_l(w^i)$
- If E is convex in w and η is small enough SGD will converge
 - very small η slow convergence
 - large η : probably it will not converge
- Scheduling for η is important (maybe decrease η with iterations)
- Single point gradients are not necessarily equal to zero as the full gradient may be. Thus this also helps to escape local minima.


Stochastic Gradient Descent - Algorithm

- SGD steps:
 - Initialize w^0 , choose η
 - Iterate over datapoints: $w^{i+1} = w^i - \eta \nabla \sum_{l=1}^L E_l(w^i)$
 - If E is convex in w and η is small enough SGD will converge
 - very small η slow convergence
 - large η : probably it will not converge
 - Scheduling for η is important (maybe decrease η with iterations)
 - Single point gradients are not necessarily equal to zero as the full gradient may be. Thus this also helps to escape local minima.
 - Initialization plays crucial role!
 - maybe perform multiple initializations
-


Stochastic gradient descent




Stochastic gradient descent




Stochastic gradient descent




Stochastic gradient descent




Stochastic gradient descent



Stochastic gradient descent



Stochastic gradient descent




Stochastic gradient descent


full gradient

$$\begin{aligned} \nabla E(\mathbf{w}^0) \\ \nabla E(\mathbf{w}^1) \\ \nabla E(\mathbf{w}^2) \end{aligned}$$

$$\begin{aligned} w^0 \\ w^1 \\ w^2 \\ w^i \end{aligned}$$



Stochastic gradient descent




Stochastic gradient descent

full gradient

gradient approximation



w_{opt}



Stochastic gradient descent

full gradient


gradient approximation



Stochastic gradient descent

full gradient

gradient approximation





w_{opt}

Stochastic gradient descent

full gradient




gradient approximation



Stochastic gradient descent

full gradient




gradient approximation



Stochastic gradient descent

full gradient

gradient approximation



Perceptron – Training : SGD

- Perceptron criterion:


$$E(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^t \mathbf{x}_n t_n = - \sum_{n \in \mathcal{M}} E_n(\mathbf{w})$$

- Sequential learning:

- Single sample SGD:

$$\mathbf{w}^{i+1} = \mathbf{w}^i - \eta \nabla E_n(\mathbf{w}^i)$$

$$\mathbf{w}^{i+1} = \mathbf{w}^i + \eta \mathbf{x}_n t_n$$



Bishop

Least Mean Squares (LMS) algorithm

- If we replace the perceptron criterion with a least squares one we get:

$$E_{LS}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^t \mathbf{x})^2$$

- If we estimate the corresponding derivative:

$$\frac{\partial E_{LS}(\mathbf{w})}{\partial \mathbf{w}} = \sum_{n=1}^N \frac{\partial \frac{1}{2} (t_n - \mathbf{w}^t \mathbf{x})^2}{\partial \mathbf{w}} = \sum_{n=1}^N -(t_n - \mathbf{w}^t \mathbf{x}) \mathbf{x}^t$$

- We get the respective Sequential Learning algorithm (SGD):

$$\mathbf{w}^{i+1} = \mathbf{w}^i + \eta \left(t_n - (\mathbf{w}^i)^t \mathbf{x} \right) \mathbf{x}^t$$

LMS algorithm - Notes

$$\mathbf{w}^{i+1} = \mathbf{w}^i + \eta \left(t_n - (\mathbf{w}^i)^t \mathbf{x} \right) \mathbf{x}^t$$

- LMS is a two-class algorithm as perceptron is.

LMS algorithm - Notes

$$\mathbf{w}^{i+1} = \mathbf{w}^i + \eta \left(t_n - (\mathbf{w}^i)^t \mathbf{x} \right) \mathbf{x}^t$$

- LMS is a two-class algorithm as perceptron is.
- Remember that for the least squares solution that we previously investigated (multiclass problem) we got:

$$\mathbf{W}_{LS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{T} = \mathbf{X}^\dagger \mathbf{T}$$

LMS algorithm - Notes

$$\mathbf{w}^{i+1} = \mathbf{w}^i + \eta \left(t_n - (\mathbf{w}^i)^t \mathbf{x} \right) \mathbf{x}^t$$

- LMS is a two-class algorithm as perceptron is.
- Remember that for the least squares solution that we previously investigated (multiclass problem) we got:

$$\mathbf{W}_{LS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{T} = \mathbf{X}^\dagger \mathbf{T}$$

- Matrix inverse is a computationally difficult to get. ($\mathcal{O}(N^3)$).
- With LMS we can overcome this problem and process data in minibatches or on a single sample base (not all data are necessary to be available at once!)

LMS algorithm - Notes

$$\mathbf{w}^{i+1} = \mathbf{w}^i + \eta \left(t_n - (\mathbf{w}^i)^t \mathbf{x} \right) \mathbf{x}^t$$


- LMS is a two-class algorithm as perceptron is.
- Remember that for the least squares solution that we previously investigated (multiclass problem) we got:

$$\mathbf{W}_{LS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{T} = \mathbf{X}^\dagger \mathbf{T}$$

- Matrix inverse is a computationally difficult to get. ($\mathcal{O}(N^3)$).
 - With LMS we can overcome this problem and process data in minibatches or on a single sample base (not all data are necessary to be available at once!)
 - LMS algorithm is also called Widrow-Hoff algorithm.
-

Neuron



$$y(x) = \mathbf{w}^t \mathbf{x} + w_0$$



Neuron

$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x} + w_0$$



$$y(\mathbf{x}) = \mathbf{w}^t \mathbf{x}$$



Neuron

$$y(x) = \mathbf{w}^t \mathbf{x}$$



- In essence, perceptron and LMS algorithms are two different approaches for the training of the model of a neuron.
- Perceptron: Rosenblatt (trained with perceptron algorithm).
- Adaline (adaptive linear element): Widrow and Hoff (trained with LMS).



Neuron

$$y(x) = \mathbf{w}^t \mathbf{x}$$

- In essence, perceptron and LMS algorithms are two different approaches for the training of the model of a neuron.
- Perceptron: Rosenblatt (trained with perceptron algorithm).
- Adaline (adaptive linear element): Widrow and Hoff (trained with LMS).
- This model is the fundamental module of Neural Networks using different activation functions.




Classification with basis functions

- So far we have considered models that work directly with the **original input** vector \mathbf{x} , e.g., $y(\mathbf{x}) = \mathbf{w}^t \mathbf{x}$.
- Nevertheless, all of the algorithms are also applicable when we first non-linearly transform the input using a vector of **basis functions** $\boldsymbol{\varphi}(\mathbf{x})$, e.g.:

$$y(\mathbf{x}) = \mathbf{w}^t \boldsymbol{\varphi}(\mathbf{x})$$

- With this trick classes that are linearly separable in the feature space $\boldsymbol{\varphi}(\mathbf{x})$ do not need to be linearly separable in the observation space \mathbf{x} .
- **Note:** Basis function are fixed, **not learned!** They need to be selected carefully. They correspond also to sources of inductive bias.


Classification with basis functions - Example



Left: original input space (x_1, x_2)

$$\varphi_1(\mathbf{x}) = \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_1)^T(\mathbf{x} - \boldsymbol{\mu}_1)\right) \text{ and } \varphi_2(\mathbf{x}) = \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_2)^T(\mathbf{x} - \boldsymbol{\mu}_2)\right)$$

Classification with basis functions - Example



Left: original input space (x_1, x_2) , **Right:** space of two gaussian basis functions with centers shown by the green crosses:

$$\varphi_1(x) = \exp\left(-\frac{1}{2}(x - \mu_1)^T(x - \mu_1)\right) \text{ and } \varphi_2(x) = \exp\left(-\frac{1}{2}(x - \mu_2)^T(x - \mu_2)\right)$$

Classification strategies (revisited)

- Probabilistic Generative Models
 - class-conditional densities $p(\mathbf{x}|\omega_i)$
 - prior class probabilities $P(\omega_i)$
 - Via Bayes we get posterior class probabilities $P(\omega_i|\mathbf{x})$
- Linear Discriminant functions
 - Direct mapping of input to target $t = y(\mathbf{x}, \mathbf{w})$
- Probabilistic Discriminative Models
 - Determining directly the posterior class probabilities $P(\omega_i|\mathbf{x})$
- Logistic Regression


Probabilistic Generative Model (revisited)

- For two classes:

$$P(\omega_1|x) = \frac{p(x|\omega_1)P(\omega_1)}{p(x)} = \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_1)P(\omega_1) + p(x|\omega_2)P(\omega_2)}$$

we can rewrite this equation as:

$$P(\omega_1|x) = \frac{1}{1+\exp(-a)} = \sigma(a)$$



where $a = \ln \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_2)P(\omega_2)}$ and $\sigma(a)$ is the *logistic sigmoid* function.


Probabilistic Generative Model (revisited)

- For number of classes $c > 2$:

$$P(\omega_i|x) = \frac{p(x|\omega_i)P(\omega_i)}{p(x)} = \frac{p(x|\omega_i)P(\omega_i)}{\sum_j p(x|\omega_j)P(\omega_j)}$$

we can rewrite this equation as:

$$P(\omega_i|x) = \frac{\exp(a_i)}{\sum_j \exp(a_j)}$$



where $a_i = \ln(p(x|\omega_i)P(\omega_i))$. This is known as the *normalized exponential* or the *softmax function*.

Probabilistic Generative Model (revisited)

- Recall that for, e.g., a two-class problem:

$$P(\omega_1|x) = \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_1)P(\omega_1) + p(x|\omega_2)P(\omega_2)} = \sigma\left(\ln \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_2)P(\omega_2)}\right)$$

Probabilistic Generative Model (revisited)

- Recall that for, e.g., a two-class problem:

$$P(\omega_1|x) = \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_1)P(\omega_1) + p(x|\omega_2)P(\omega_2)} = \sigma\left(\ln \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_2)P(\omega_2)}\right)$$

$$g(\mathbf{x}) = g_1(\mathbf{x}) - g_2(\mathbf{x})$$

$$g(\mathbf{x}) = P(\omega_1 | \mathbf{x}) - P(\omega_2 | \mathbf{x})$$

$$g(\mathbf{x}) = \ln \frac{p(\mathbf{x}|\omega_1)}{p(\mathbf{x}|\omega_2)} + \ln \frac{P(\omega_1)}{P(\omega_2)}$$

slide 38 of Bayes Theory lecture

Probabilistic Generative Model (revisited)

- Recall that for, e.g., a two-class problem:

$$P(\omega_1|x) = \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_1)P(\omega_1) + p(x|\omega_2)P(\omega_2)} = \sigma\left(\ln \frac{p(x|\omega_1)P(\omega_1)}{p(x|\omega_2)P(\omega_2)}\right)$$

and if we adopt Gaussian class conditional densities with same covariance matrix across classes we end up with:

$$g(\mathbf{x}) = g_1(\mathbf{x}) - g_2(\mathbf{x})$$

$$P(\omega_1|x) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

$$g(\mathbf{x}) = P(\omega_1 | \mathbf{x}) - P(\omega_2 | \mathbf{x})$$

where

$$\mathbf{w} = \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2) \text{ and}$$

$$g(\mathbf{x}) = \ln \frac{p(\mathbf{x}|\omega_1)}{p(\mathbf{x}|\omega_2)} + \ln \frac{P(\omega_1)}{P(\omega_2)}$$

slide 38 of Bayes Theory lecture

$$w_0 = -\frac{1}{2}\boldsymbol{\mu}_1\boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_1 + \frac{1}{2}\boldsymbol{\mu}_2\boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_2 + \ln \frac{P(\omega_1)}{P(\omega_2)}$$

- We can use Maximum Likelihood to estimate $\boldsymbol{\mu}_1$, $\boldsymbol{\mu}_2$ and $\boldsymbol{\Sigma}$.

Logistic Regression

- Thus we have shown that the posterior class probability **under certain conditions** can be written as a logistic sigmoid acting on a linear function of the feature vector:

$$P(\omega_1|x) = y(x) = \sigma(w^T x)$$

with $P(\omega_2|x) = 1 - P(\omega_1|x)$.

- Using basis functions we can rewrite this as:

$$P(\omega_1|\varphi(x)) = y(\varphi(x)) = \sigma(w^T \varphi(x))$$

with $P(\omega_2|\varphi(x)) = 1 - P(\omega_1|\varphi(x))$.

- In statistics this model is known as *logistic regression*. However, this is a model for **classification not regression**.

Logistic Regression

$$P(\omega_1 | \boldsymbol{\varphi}(\mathbf{x})) = y(\boldsymbol{\varphi}(\mathbf{x})) = \sigma(\mathbf{w}^T \boldsymbol{\varphi}(\mathbf{x}))$$

with $P(\omega_2 | \boldsymbol{\varphi}(\mathbf{x})) = 1 - P(\omega_1 | \boldsymbol{\varphi}(\mathbf{x}))$.

- For a d –dimensional feature space $\boldsymbol{\varphi}$, this model has **$d + 1$ adjustable parameters** (weight of the linear function).
- If we have tried to fit gaussian class conditional densities we would need to estimate $2d$ parameters for the means and $\mathcal{O}(d^2)$ parameters for the shared covariance matrix!
- Thus, even if the logistic regression parameters scale **linearly with d** the probabilistic generative model approach scales **quadratically with d** !
- There is clear advantage to work with the logistic regression model directly!

Logistic Regression–Training (Maximum Likelihood)

$$P(\omega_1 | \boldsymbol{\varphi}(\mathbf{x})) = y(\boldsymbol{\varphi}(\mathbf{x})) = \sigma(\mathbf{w}^T \boldsymbol{\varphi}(\mathbf{x}))$$

- In order to adjust the weights of the logistic regression model we can try Maximum Likelihood.
- We will make use of the fact that:

$$\frac{d\sigma(a)}{da} = \sigma(a)(1 - \sigma(a))$$

- Thus with a training set of N samples of the form $\{\boldsymbol{\varphi}(\mathbf{x}_n), t_n\}$ with $t_n \in \{0,1\}$ the likelihood function can be written as:

$$p(\mathbf{t}|\mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n},$$

where $y_n = P(\omega_1 | \boldsymbol{\varphi}(\mathbf{x}_n)) = \sigma(\mathbf{w}^T \boldsymbol{\varphi}(\mathbf{x}_n))$ and $\mathbf{t} = (t_1, \dots, t_N)$

Logistic Regression–Training (Maximum Likelihood)

- We can **define an error function** and instead of maximizing the likelihood:

$$p(\mathbf{t}|\mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n},$$

we can **minimize the error function** which is the negative logarithm of the likelihood:

$$E(\mathbf{w}) = -\ln p(\mathbf{t}|\mathbf{w}) = -\sum_{n=1}^N [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]$$

- This error function is called *cross entropy loss* and its gradient is:

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \boldsymbol{\varphi}(\mathbf{x}_n)$$

Logistic Regression–Training(SGD)

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \boldsymbol{\varphi}(x_n) = \sum_{n=1}^N \nabla E_n(\mathbf{w})$$

- Unfortunately, we cannot derive a closed form solution for \mathbf{w} as $y(\boldsymbol{\varphi}(x)) = \sigma(\mathbf{w}^T \boldsymbol{\varphi}(x))$ is nonlinear on \mathbf{w} .
- Nevertheless, $E(\mathbf{w})$ is a convex function of \mathbf{w} and we could use a sequential learning algorithm for the logistic regression model training.
- In particular with SGD we get:

$$\mathbf{w}^{i+1} = \mathbf{w}^i - \eta \nabla E_n(\mathbf{w}^i) \quad \rightarrow \quad \mathbf{w}^{i+1} = \mathbf{w}^i - \eta (y_n - t_n) \boldsymbol{\varphi}(x_n)$$

- **Classification:** a test sample is assigned to ω_1 if $P(\omega_1 | \boldsymbol{\varphi}(x)) = \sigma(\mathbf{w}^T \boldsymbol{\varphi}(x)) \geq \frac{1}{2}$



ARISTOTLE UNIVERSITY OF THESSALONIKI



FACULTY OF ENGINEERING

Questions?

Pattern Recognition & Machine Learning
Linear Discriminant Functions and Models