



ARISTOTLE UNIVERSITY OF THESSALONIKI



FACULTY OF ENGINEERING

# Pattern Recognition & Machine Learning

## *Model Combination - Trees*

**Panagiotis C. Petrantonakis**

*Assistant Professor*

Dept. of Electrical and Computer Engineering

[ppetrant@ece.auth.gr](mailto:ppetrant@ece.auth.gr)

Fall Semester

# Until now...

---

- We have talked about many different models for classification.
- It is often the case that the overall performance could be improved by model combination!
- Example: You have **multiple NN models** trained on **different datasets** for the same regression problem. For a new, test sample, the **average of the predictions** made by each model would probably lead to a better result than the outcome of a single model.
- The group of models that are combined like this are called *committees*!
  - Why *committees* lead to better results?

# Test error of a classifier

---

- You can think about test error in terms of a decomposition into **two terms**.

# Test error of a classifier

---

- You can think about test error in terms of a decomposition into **two terms**.
- Let  $f$  be a learned classifier, selected from a set  $\mathcal{F}$ , i.e., all possible classifiers using a fixed representation, e.g., trees with certain depth, NNs with certain number of neurons in the hidden layer etc.

# Test error of a classifier

---

- You can think about test error in terms of a decomposition into **two terms**.
- Let  $f$  be a learned classifier, selected from a set  $\mathcal{F}$ , i.e., all possible classifiers using a fixed representation, e.g., trees with certain depth, NNs with certain number of neurons in the hidden layer etc.
- Then the error of the classifier on a test set can be decomposed as:

$$\begin{aligned} \text{error}(f) &= \mathbb{E}_{\mathcal{D}} \left[ (f(\mathbf{x}; \mathcal{D}) - h(\mathbf{x}))^2 \right] = \\ &= (\mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] - h(\mathbf{x}))^2 + \mathbb{E}_{\mathcal{D}}[(f(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})])^2] \end{aligned}$$

where  $\mathcal{D}$  is a particular dataset and  $h(\mathbf{x})$  is the function to be approximated.

---

# Bias/Variance Trade-off

---

$$\mathbb{E}_{\mathcal{D}} \left[ (f(\mathbf{x}; \mathcal{D}) - h(\mathbf{x}))^2 \right] = \underbrace{\left( \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] - h(\mathbf{x}) \right)^2}_{\text{Bias}} + \underbrace{\mathbb{E}_{\mathcal{D}}[(f(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})])^2]}_{\text{Variance}}$$

- **Bias:** measures the quality of the model family (e.g., NNs with certain number of neurons in hidden layer).
  - Suppose someone gave you infinite data to train your model; How well would you do with this particular **model family**?
- **Variance:** measures how far the actual, learned classifier (train over limited data) is from the optimal classifier of the family, i.e., the one trained with infinite data.
  - In essence, this tells you **how much you have to pay** due to the fact that you do not have infinite amount of data!
  - **Sensitivity** of the model to the individual dataset it was trained on.

# Bias/Variance Trade-off

---

$$\mathbb{E}_{\mathcal{D}} \left[ (f(\mathbf{x}; \mathcal{D}) - h(\mathbf{x}))^2 \right] = \underbrace{\left( \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] - h(\mathbf{x}) \right)^2}_{\text{Bias}} + \underbrace{\mathbb{E}_{\mathcal{D}}[(f(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})])^2]}_{\text{Variance}}$$

- There is a fundamental **trade-off** between Bias and Variance in practical applications.
- As you make your model family more complex, e.g., **complex NN architecture**, you make  $\mathcal{F}$  bigger. This, will cause a **decrease in bias** term as the approximation with infinite data will be much better.
- On the other hand this will also cause and **increase in the variance** term as more complex representation with limited data will most probably lead to overfitting!

# Bias/Variance Trade-off - Example

- Assume we have  $L$  datasets of  $N$  points to approximate the function:

$$h(x) = \sin(2\pi x) + \varepsilon$$

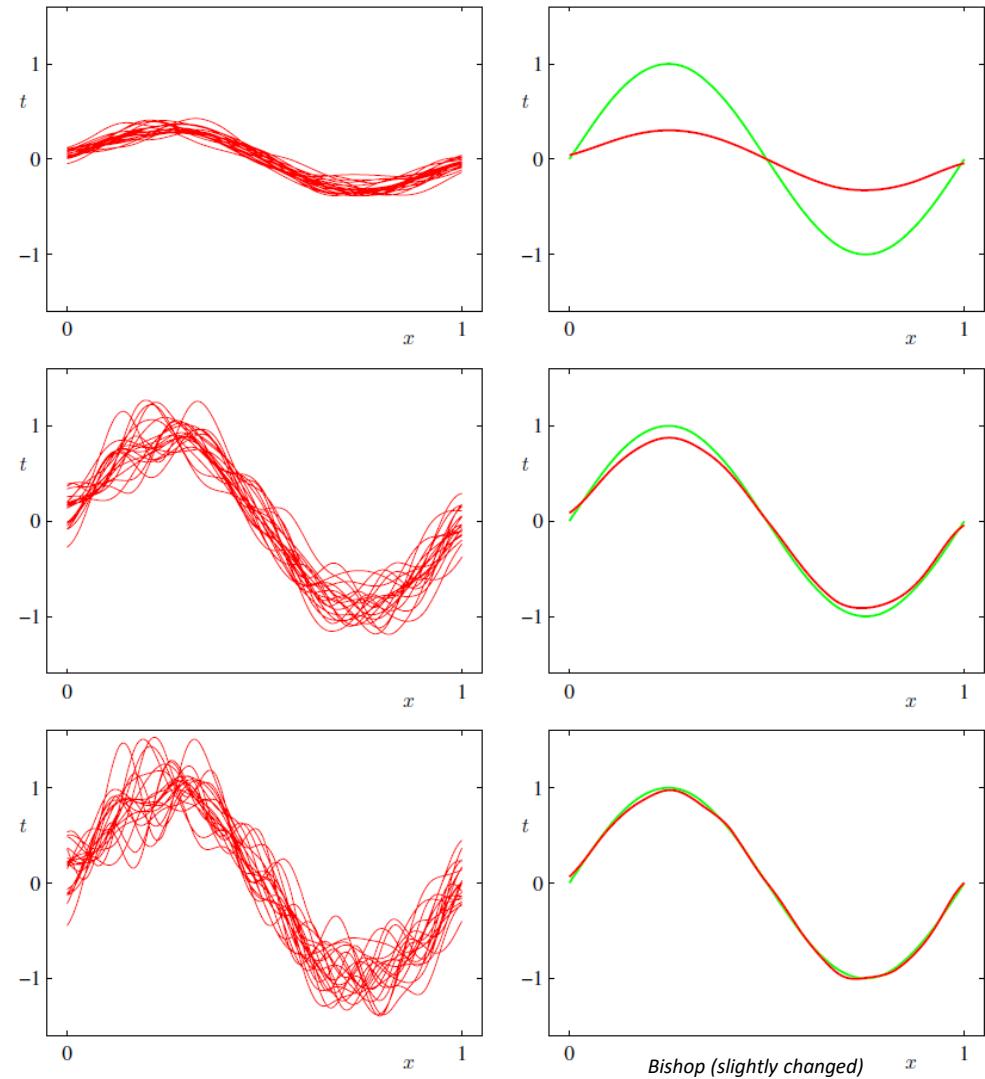
The final outcome of the committee will be:

$$\mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] = \frac{1}{L} \sum_{l=1}^L f^l(x)$$

first row: high bias, low variance

second row: lower bias, high variance

third row: low bias, higher variance



# Averaging predictions

---

- Averaging predictions of models **trained on different datasets** leads to reduction of the variance!
  - Low bias/high variance leads to low variance/low bias after averaging
- Thus, when we average a set of low-bias (high variance) models, i.e., complex models, we obtain accurate predictions. **That's why committees work!**
- Pitfall: we only have **one** single dataset!
- How could we artificially introduce variability between different models within a committee?

# Averaging predictions

---

- Averaging predictions of models trained on different datasets leads to reduction of the variance!
  - Low bias/high variance leads to low variance/low bias after averaging
- Thus, when we average a set of low-bias (high variance) models, i.e., complex models, we obtain accurate predictions. **That's why committees work!**
- Pitfall: we only have **one** single dataset!
- How could we artificially introduce variability between different models within a committee?
  - Bootstrap datasets!

# Bootstrapping Datasets

---

- Suppose we have the dataset  $X = \{x_1, \dots, x_N\}$
- we can create  $B$  new datasets by simply drawing  $N$  points at random from  $X$  with replacement!
  - Some samples may appear multiple times in a new dataset after bootstrapping.
  - Some samples of the initial dataset may not appear at all.
- Thus, I get  $B$  new datasets:  $\{X_1, \dots, X_B\}$ .

# Bootstrapping Datasets

---

- Suppose we have the dataset  $X = \{x_1, \dots, x_N\}$
- we can create  $B$  new datasets by simply drawing  $N$  points at random from  $X$  with replacement!
  - Some samples may appear multiple times in a new dataset after bootstrapping.
  - Some samples of the initial dataset may not appear at all.
- Thus, I get  $B$  new datasets:  $\{X_1, \dots, X_B\}$ .
- Now we can use these  $B$  training sets to train  $B$  different models (committee members).
- The final decision of the committee will be:  $f_{com} = \frac{1}{B} \sum_{b=1}^B f_b(x)$ 
  - Bootstrap aggregation/bagging

# Regression with bootstrapping

---

- Suppose the ground truth function that we need to predict is  $h(\mathbf{x})$ .
- Each model gives a prediction:  $f_b(\mathbf{x}) = h(\mathbf{x}) + \epsilon_b(\mathbf{x})$ 
  - error of the model:  $\epsilon_b(\mathbf{x})$
- The average sum-of-squares error of model  $b$  is :

$$\mathbb{E}_{\mathbf{x}} \left[ (f_b(\mathbf{x}) - h(\mathbf{x}))^2 \right] = \mathbb{E}_{\mathbf{x}} [\epsilon_b(\mathbf{x})^2]$$

- The average error of  $B$  individual models is:

$$E_{ave} = \frac{1}{B} \sum_{b=1}^B \mathbb{E}_{\mathbf{x}} [\epsilon_b(\mathbf{x})^2]$$

# Regression with bootstrapping

---

- The expected error of the committee is:

$$E_{com} = \mathbb{E}_{\mathbf{x}} \left[ (f_{com}(\mathbf{x}) - h(\mathbf{x}))^2 \right] = \mathbb{E}_{\mathbf{x}} \left[ \left( \frac{1}{B} \sum_{b=1}^B f_b(\mathbf{x}) - h(\mathbf{x}) \right)^2 \right]$$

# Regression with bootstrapping

---

- The expected error of the committee is:

$$E_{com} = \mathbb{E}_{\mathbf{x}} \left[ (f_{com}(\mathbf{x}) - h(\mathbf{x}))^2 \right] = \mathbb{E}_{\mathbf{x}} \left[ \left( \frac{1}{B} \sum_{b=1}^B f_b(\mathbf{x}) - h(\mathbf{x}) \right)^2 \right]$$

- If we assume uncorrelated error between the different models then we can prove that:

$$E_{com} = \frac{1}{B} E_{ave}$$

# Regression with bootstrapping

---

- The expected error of the committee is:

$$E_{com} = \mathbb{E}_{\mathbf{x}} \left[ (f_{com}(\mathbf{x}) - h(\mathbf{x}))^2 \right] = \mathbb{E}_{\mathbf{x}} \left[ \left( \frac{1}{B} \sum_{b=1}^B f_b(\mathbf{x}) - h(\mathbf{x}) \right)^2 \right]$$

- If we assume uncorrelated error between the different models then we can prove that:

$$E_{com} = \frac{1}{B} E_{ave}$$

- Nevertheless, the way we defined the datasets after bootstrapping, it is not possible that the errors are not correlated.

# Regression with bootstrapping

---

- The expected error of the committee is:

$$E_{com} = \mathbb{E}_{\mathbf{x}} \left[ (f_{com}(\mathbf{x}) - h(\mathbf{x}))^2 \right] = \mathbb{E}_{\mathbf{x}} \left[ \left( \frac{1}{B} \sum_{b=1}^B f_b(\mathbf{x}) - h(\mathbf{x}) \right)^2 \right]$$

- If we assume uncorrelated error between the different models then we can prove that:

$$E_{com} = \frac{1}{B} E_{ave}$$

- Nevertheless, the way we defined the datasets after bootstrapping, it is not possible that the errors are not correlated.
- Even in the correlated case, it can be proved that:  $E_{com} \leq E_{ave}$

# Regression with bootstrapping

---

- The expected error of the committee is:

$$E_{com} = \mathbb{E}_{\mathbf{x}} \left[ (f_{com}(\mathbf{x}) - h(\mathbf{x}))^2 \right] = \mathbb{E}_{\mathbf{x}} \left[ \left( \frac{1}{B} \sum_{b=1}^B f_b(\mathbf{x}) - h(\mathbf{x}) \right)^2 \right]$$

- If we assume uncorrelated error between the different models then we can prove that:

$$E_{com} = \frac{1}{B} E_{ave}$$

- Nevertheless, the way we defined the datasets after bootstrapping, it is not possible that the errors are not correlated.
  - Even in the correlated case, it can be proved that:  $E_{com} \leq E_{ave}$
  - **Basic strategy:** choose  $B$  models with low bias (**complex models that can overfit**). Then bootstrap aggregated error will be lower than the average error of the individual models.
-

# Feature Bagging

---

- Feature Bagging: create new datasets by sampling **subsets of features** of length  $M < D$  for each learner.
  - Also called ‘random subspace method’

# Feature Bagging

---

- Feature Bagging: create new datasets by sampling **subsets of features** of length  $M < D$  for each learner.
  - Also called ‘random subspace method’
- Example:
  - Assume dataset  $X = \{x_1, \dots, x_N\}$  with samples  $x_n = \{x_1, x_2, x_3, x_4, x_5, x_6\}$ . I create new datasets  $Y^1 = \{y_1^1, \dots, y_N^1\}$  with  $y_n^1 = \{x_1, x_2, x_6\}$  and  $Y^2 = \{y_1^2, \dots, y_N^2\}$  with  $y_n^2 = \{x_3, x_4, x_6\}$ .
  - Train different models with different datasets after feature bagging

# Feature Bagging

---

- Feature Bagging: create new datasets by sampling **subsets of features** of length  $M < D$  for each learner.
    - Also called ‘random subspace method’
  - Example:
    - Assume dataset  $X = \{x_1, \dots, x_N\}$  with samples  $x_n = \{x_1, x_2, x_3, x_4, x_5, x_6\}$ . I create new datasets  $Y^1 = \{y_1^1, \dots, y_N^1\}$  with  $y_n^1 = \{x_1, x_2, x_6\}$  and  $Y^2 = \{y_1^2, \dots, y_N^2\}$  with  $y_n^2 = \{x_3, x_4, x_6\}$ .
    - Train different models with different datasets after feature bagging
  - It works well if features are **uncorrelated**.
  - **Advantage:** pushes learners to **not** over focus on features that predict well on certain training set but do not generalize well to new data!
  - Works well when I have too many features and too few samples.
-

# Boosting

---

- Bootstrapping and Feature Bagging methods create new datasets and train **in parallel** different models. The final decision was made by averaging.

# Boosting

---

- Bootstrapping and Feature Bagging methods create new datasets and train **in parallel** different models. The final decision was made by averaging.
- *Boosting* algorithms use initial dataset but train, **sequentially**, different models.

# Boosting

---

- Bootstrapping and Feature Bagging methods create new datasets and train **in parallel** different models. The final decision was made by averaging.
- *Boosting* algorithms use initial dataset but train, **sequentially**, different models.
- In boosting, we use **simple classifiers** (weak learners). Each one of these classifiers is trained using a weighted form of the dataset.

# Boosting

---

- Bootstrapping and Feature Bagging methods create new datasets and train **in parallel** different models. The final decision was made by averaging.
- *Boosting* algorithms use initial dataset but train, **sequentially**, different models.
- In boosting, we use **simple classifiers** (weak learners). Each one of these classifiers is trained using a weighted form of the dataset.
- The weighting coefficient associated with each datapoint of the dataset, depends on the performance of the previous model.

# Boosting

---

- Bootstrapping and Feature Bagging methods create new datasets and train **in parallel** different models. The final decision was made by averaging.
- *Boosting* algorithms use initial dataset but train, **sequentially**, different models.
- In boosting, we use **simple classifiers** (weak learners). Each one of these classifiers is trained using a weighted form of the dataset.
- The weighting coefficient associated with each datapoint of the dataset, depends on the performance of the previous model.
- Samples that are misclassified are given greater weight for the subsequent classifier in the sequence.

# Boosting

---

- Bootstrapping and Feature Bagging methods create new datasets and train **in parallel** different models. The final decision was made by averaging.
- *Boosting* algorithms use initial dataset but train, **sequentially**, different models.
- In boosting, we use **simple classifiers** (weak learners). Each one of these classifiers is trained using a weighted form of the dataset.
- The weighting coefficient associated with each datapoint of the dataset, depends on the performance of the previous model.
- Samples that are misclassified are given greater weight for the subsequent classifier in the sequence.
- Thus,
  - **Bootstrap-bagging:** Complex model with low bias. Decrease variance by averaging.
  - **Boosting:** Simple models with high bias (underfit). Decrease bias (and variance) by focusing on error of previous models.

# Boosting

---

- Bootstrapping and Feature Bagging methods create new datasets and train **in parallel** different models. The final decision was made by averaging.
  - *Boosting* algorithms use initial dataset but train, **sequentially**, different models.
  - In boosting, we use **simple classifiers** (weak learners). Each one of these classifiers is trained using a weighted form of the dataset.
  - The weighting coefficient associated with each datapoint of the dataset, depends on the performance of the previous model.
  - Samples that are misclassified are given greater weight for the subsequent classifier in the sequence.
  - Thus,
    - **Bootstrap-bagging:** Complex model with low bias. Decrease variance by averaging.
    - **Boosting:** Simple models with high bias (underfit). Decrease bias (and variance) by focusing on error of previous models.
  - Boosting, final decision:  $Y_{com}(\mathbf{x}) = \text{sign}(\sum_{m=1}^M a_m y_m(\mathbf{x}))$ ,  $y_m(\mathbf{x}) \in \{-1, 1\}$
-

# Boosting: Binary Classification

---

- Assume  $X = \{(x_n, t_n)\}_{n=1}^N$  where  $t_n \in \{-1, 1\}$

# Boosting: Binary Classification

---

- Assume  $X = \{(\mathbf{x}_n, t_n)\}_{n=1}^N$  where  $t_n \in \{-1, 1\}$
- Each data sample has an associated weight  $w_n$ . Initialization:  $w_n = \frac{1}{N}$
- Assume that our base classifier produces labels  $y_m(\mathbf{x}) \in \{-1, 1\}$

# Boosting: Binary Classification

---

- Assume  $X = \{(\mathbf{x}_n, t_n)\}_{n=1}^N$  where  $t_n \in \{-1, 1\}$
- Each data sample has an associated weight  $w_n$ . Initialization:  $w_n = \frac{1}{N}$
- Assume that our base classifier produces labels  $y_m(\mathbf{x}) \in \{-1, 1\}$
- The most popular boosting algorithm is AdaBoost (Adaptive Boosting)
  - At each step a new classifier is trained on weighted dataset
  - Weights of misclassified data samples increase.
  - Committee decision is formed by weighted base classifiers.

# AdaBoost Algorithm

---

1. Initialize weights:  $w_n^1 = \frac{1}{N}, n = 1, \dots, N$

# AdaBoost Algorithm

---

1. Initialize weights:  $w_n^1 = \frac{1}{N}, n = 1, \dots, N$
2. for  $m = 1, \dots, M$ :
  - a) Train  $y_m(\mathbf{x})$  to minimize  $J_m = \sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]$
  - b) Compute weighted error rates  $\varepsilon_m = \frac{\sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]}{\sum_{n=1}^N w_n^m}$   
and model weight  $a_m = \ln\left(\frac{1-\varepsilon_m}{\varepsilon_m}\right)$
  - a) Update weights  $w_n^{m+1} = w_n^m \exp\{a_m I[y_m(\mathbf{x}_n) \neq t_n]\}$

# AdaBoost Algorithm

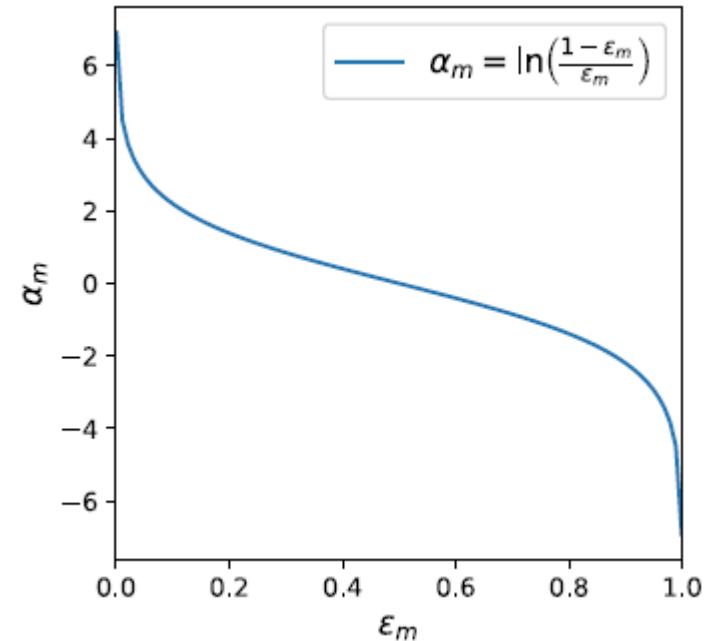
---

1. Initialize weights:  $w_n^1 = \frac{1}{N}, n = 1, \dots, N$
  2. for  $m = 1, \dots, M$ :
    - a) Train  $y_m(\mathbf{x})$  to minimize  $J_m = \sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]$
    - b) Compute weighted error rates  $\varepsilon_m = \frac{\sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]}{\sum_{n=1}^N w_n^m}$   
and model weight  $a_m = \ln\left(\frac{1-\varepsilon_m}{\varepsilon_m}\right)$
  3. Make prediction by  $Y_M(\mathbf{x}) = \text{sign}\left(\sum_{m=1}^M a_m y_m(\mathbf{x})\right)$
-

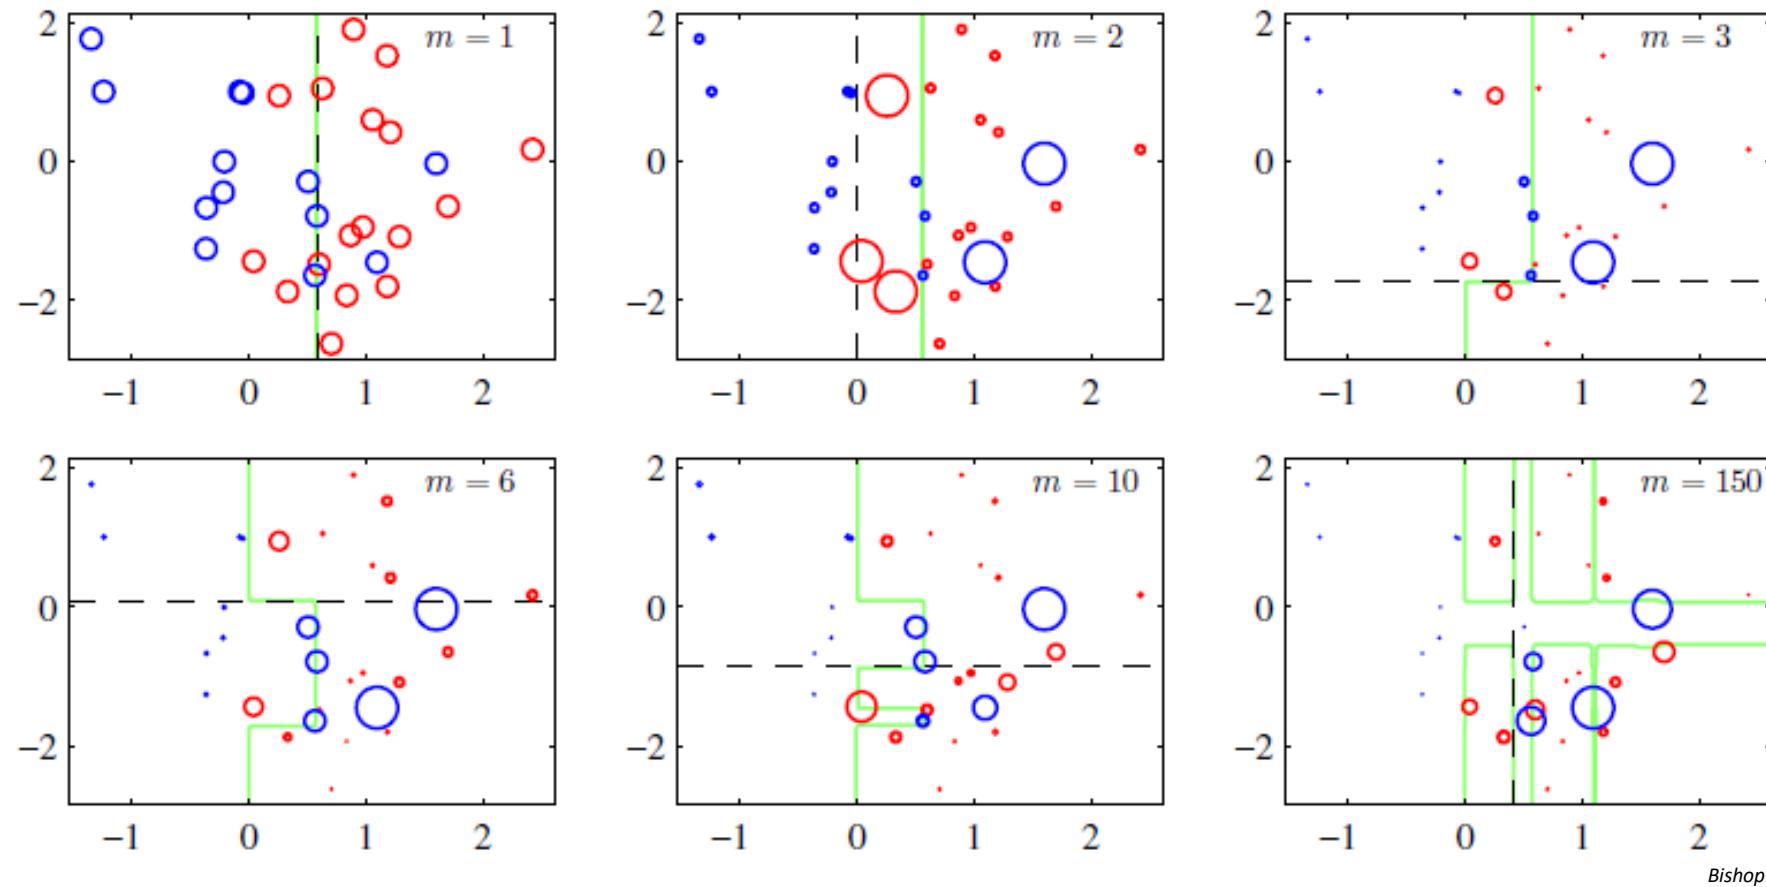
# AdaBoost Algorithm

---

1. Initialize weights:  $w_n^1 = \frac{1}{N}, n = 1, \dots, N$
  2. for  $m = 1, \dots, M$ :
    - a) Train  $y_m(\mathbf{x})$  to minimize  $J_m = \sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]$
    - b) Compute weighted error rates  $\varepsilon_m = \frac{\sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]}{\sum_{n=1}^N w_n^m}$   
and model weight  $\alpha_m = \ln\left(\frac{1-\varepsilon_m}{\varepsilon_m}\right)$
    - a) Update weights  $w_n^{m+1} = w_n^m \exp\{\alpha_m I[y_m(\mathbf{x}_n) \neq t_n]\}$
  3. Make prediction by  $Y_M(\mathbf{x}) = \text{sign}\left(\sum_{m=1}^M \alpha_m y_m(\mathbf{x})\right)$
- 



# AdaBoost Algorithm - Example



Dashed line: current base classifier boundary, green line: committee's decision boundary, cycles: training samples, radius of cycles: weight of the sample for current base classifier.

# Interpretation of AdaBoost

---

- AdaBoost is actually minimizing an exponential error function:

$$E_m = \sum_{n=1}^N \exp[-t_n f_m(\mathbf{x}_n)]$$

where  $t_n f_m(\mathbf{x}_n) > 0$  if the sample  $\mathbf{x}_n$  is correctly classified and  $t_n f_m(\mathbf{x}_n) \leq 0$  otherwise.

# Interpretation of AdaBoost

---

- AdaBoost is actually minimizing an exponential error function:

$$E_m = \sum_{n=1}^N \exp[-t_n f_m(\mathbf{x}_n)]$$

where  $t_n f_m(\mathbf{x}_n) > 0$  if the sample  $\mathbf{x}_n$  is correctly classified and  $t_n f_m(\mathbf{x}_n) \leq 0$  otherwise. Moreover:

$$f_m(\mathbf{x}_n) = \frac{1}{2} \sum_{i=1}^m a_i y_i(\mathbf{x}_n)$$

is the linear combination of base classifiers  $y_i(\mathbf{x}_n)$ , where  $m$  is the number of committee members.

# Interpretation of AdaBoost

---

- AdaBoost is actually minimizing an exponential error function:

$$E_m = \sum_{n=1}^N \exp[-t_n f_m(\mathbf{x}_n)]$$

where  $t_n f_m(\mathbf{x}_n) > 0$  if the sample  $\mathbf{x}_n$  is correctly classified and  $t_n f_m(\mathbf{x}_n) \leq 0$  otherwise. Moreover:

$$f_m(\mathbf{x}_n) = \frac{1}{2} \sum_{i=1}^m a_i y_i(\mathbf{x}_n)$$

is the linear combination of base classifiers  $y_i(\mathbf{x}_n)$ , where  $m$  is the number of committee members.

- Goal:** minimize  $E_m$  w.r.t.  $a_i$  and the parameters of  $y_i(\mathbf{x}_n)$ 
    - Sequential minimization:* fix  $y_1(\mathbf{x}), \dots, y_{m-1}(\mathbf{x})$  and  $a_1, \dots, a_{m-1}$  and minimize  $E_m$  w.r.t.  $a_m$  and the parameters of  $y_m(\mathbf{x}_n)$
-

# Derivation of AdaBoost - 1

$$f_m(\mathbf{x}_n) = \frac{1}{2} \sum_{i=1}^m a_i y_i(\mathbf{x}_n)$$

- Error function:

$$\begin{aligned} E_m &= \sum_{n=1}^N \exp\left(-t_n f_{m-1}(\mathbf{x}_n) - \frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right) \\ &= \sum_{n=1}^N w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right) \end{aligned}$$

where  $w_n^m = \exp(-t_n f_{m-1}(\mathbf{x}_n))$  is known (fixed) based on the trainings so far!

- Assume:
  - $C_m$  to be the set of samples that are correctly classified by  $y_m(\mathbf{x})$ , i.e.,  $t_n y_m(\mathbf{x}_n) = 1$
  - $M_m$  to be the set of samples that are misclassified by  $y_m(\mathbf{x})$ , i.e.,  $t_n y_m(\mathbf{x}_n) = -1$
- Then we can rewrite error function as:

# Derivation of AdaBoost - 2

---

$$\begin{aligned} E_m &= \sum_{n=1}^N w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right) = e^{-\frac{a_m}{2}} \sum_{n \in C_m} w_n^m + e^{\frac{a_m}{2}} \sum_{n \in M_m} w_n^m \\ &= \left(e^{\frac{a_m}{2}} - e^{-\frac{a_m}{2}}\right) \sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n] + e^{-\frac{a_m}{2}} \sum_{n=1}^N w_n^m \end{aligned}$$

- Thus, minimization of  $E_m$  w.r.t.  $y_m(\mathbf{x}_n)$  minimizes  $J_m = \sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]$
- minimization of  $E_m$  w.r.t.  $a_m$ :  $\frac{\partial E_m}{\partial a_m} = 0$  :

$$a_m = \ln\left(\frac{1 - \varepsilon_m}{\varepsilon_m}\right) \text{ where } \varepsilon_m = \frac{\sum_{n=1}^N w_n^m I[y_m(\mathbf{x}_n) \neq t_n]}{\sum_{n=1}^N w_n^m}$$

# Derivation of AdaBoost - 3

---

- As soon as we have found  $y_m(\mathbf{x}_n)$  and  $a_m$ , we subsequently minimize  $E_{m+1}$  w.r.t.  $y_{m+1}(\mathbf{x}_n)$  and  $a_{m+1}$ :

$$\begin{aligned} E_{m+1} &= \sum_{n=1}^N \exp[-t_n f_{m+1}(\mathbf{x}_n)] = \\ &= \sum_{n=1}^N w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right) \exp\left(-\frac{1}{2} t_n a_{m+1} y_{m+1}(\mathbf{x}_n)\right) = \\ &= \sum_{n=1}^N w_n^{m+1} \exp\left(-\frac{1}{2} t_n a_{m+1} y_{m+1}(\mathbf{x}_n)\right) \end{aligned}$$

where

$$w_n^{m+1} = w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right)$$

---

# Derivation of AdaBoost - 4

---

- Weight updates:

$$w_n^{m+1} = w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right)$$

# Derivation of AdaBoost - 4

---

- Weight updates:

$$w_n^{m+1} = w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right)$$

- If we use:  $t_n y_m(\mathbf{x}_n) = 1 - 2I[y_m(\mathbf{x}_n) \neq t_n]$

# Derivation of AdaBoost - 4

---

- Weight updates:

$$w_n^{m+1} = w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right)$$

- If we use:  $t_n y_m(\mathbf{x}_n) = 1 - 2I[y_m(\mathbf{x}_n) \neq t_n]$
- Then  $w_n^{m+1} = w_n^m \exp\left(-\frac{a_m}{2}\right) \exp(a_m I[y_m(\mathbf{x}_n) \neq t_n]) =$   
 $w_n^m \exp(a_m I[y_m(\mathbf{x}_n) \neq t_n])$
- The term  $\exp\left(-\frac{a_m}{2}\right)$  is independent of  $n$  and, thus, can be discarded.

# Derivation of AdaBoost - 4

---

- Weight updates:

$$w_n^{m+1} = w_n^m \exp\left(-\frac{1}{2} t_n a_m y_m(\mathbf{x}_n)\right)$$

- If we use:  $t_n y_m(\mathbf{x}_n) = 1 - 2I[y_m(\mathbf{x}_n) \neq t_n]$
  - Then  $w_n^{m+1} = w_n^m \exp\left(-\frac{a_m}{2}\right) \exp(a_m I[y_m(\mathbf{x}_n) \neq t_n]) =$   
 $w_n^m \exp(a_m I[y_m(\mathbf{x}_n) \neq t_n])$
  - The term  $\exp\left(-\frac{a_m}{2}\right)$  is independent of  $n$  and, thus, can be discarded.
  - Finally,  $sign(f_m(\mathbf{x}_n)) = sign\left(\frac{1}{2} \sum_{i=1}^m a_i y_i(\mathbf{x}_n)\right) = sign(\sum_{i=1}^m a_i y_i(\mathbf{x}_n))$
-

# Pros and Cons of AdaBoost

---

- The selection of the exponential error function, makes AdaBoost a very simple algorithm with easy implementation.
- It is sensitive to large negative outliers of the term  $t_n y_m(x_n)$  (in case there is not suitable activation function to get +1 or -1).
- Exponential error function can not be interpreted as a well defined probabilistic model.
- AdaBoost does not generalize easily to  $c > 2$ 
  - There are many alterations of the algorithm, though.

# Decision Trees (revisited)

# Combining models

---

- Committees
  - Bootstrap Aggregation
  - Random Subspace methods (feature bagging)
  - Boosting
- We will now revisit the Decision Tree model and will extend the approach to Random Forests.
- In essence, decisions trees with bootstrapping and random subspaces methods lead to the random forests realization.

# Decision Trees (DT)

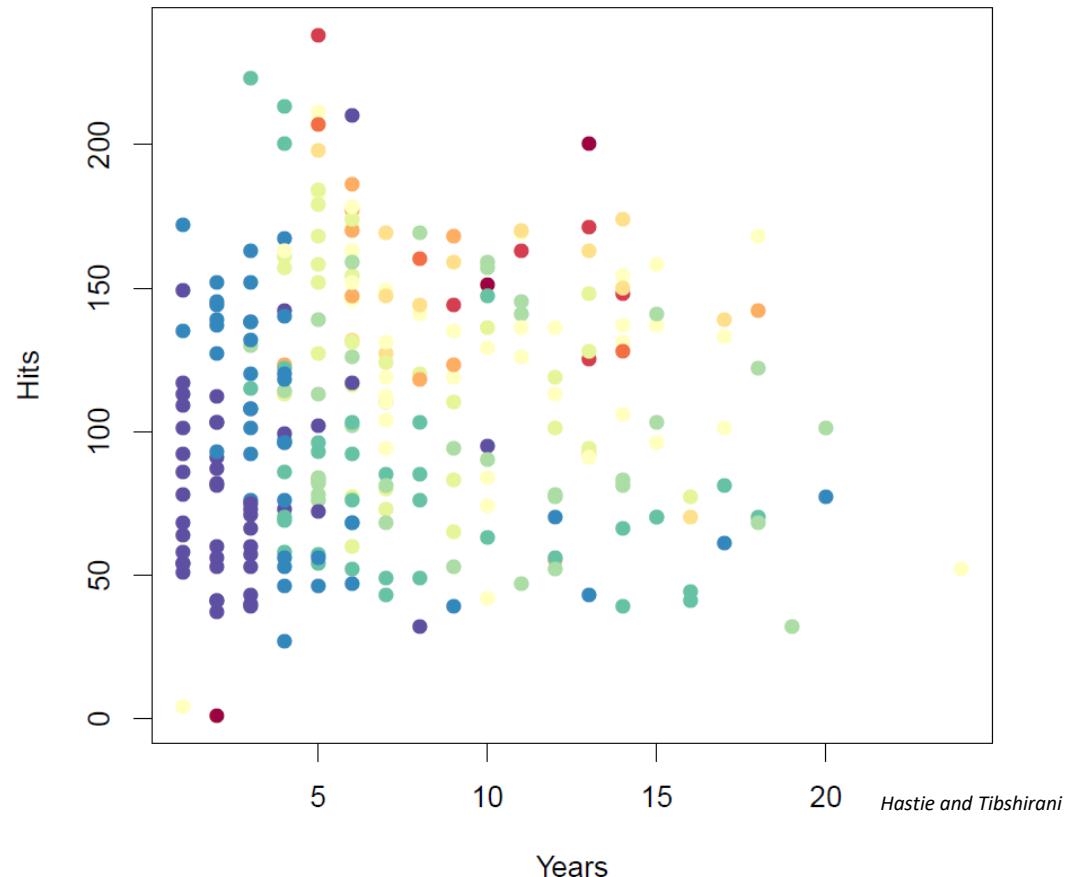
---

- DT methods can be used both for regression and classification
  - They result in stratification of segmentation of the feature space into a number of **rectangular regions**.
  - The rules that segment the feature space into distinct regions can be summarized using tree, thus the term Decision Trees.
  - DT are suitable for **intuitive interpretation** of the outcome!
  - Nevertheless, they are in general **weak classifiers** (or regressors) and cannot compete with the supervised approaches that we have seen so far.
  - Combining a large number of DT grown with bagging and/or boosting approaches can often result in great improvement in prediction accuracy at the expense of loss interpretation.
-

# DT - Regression

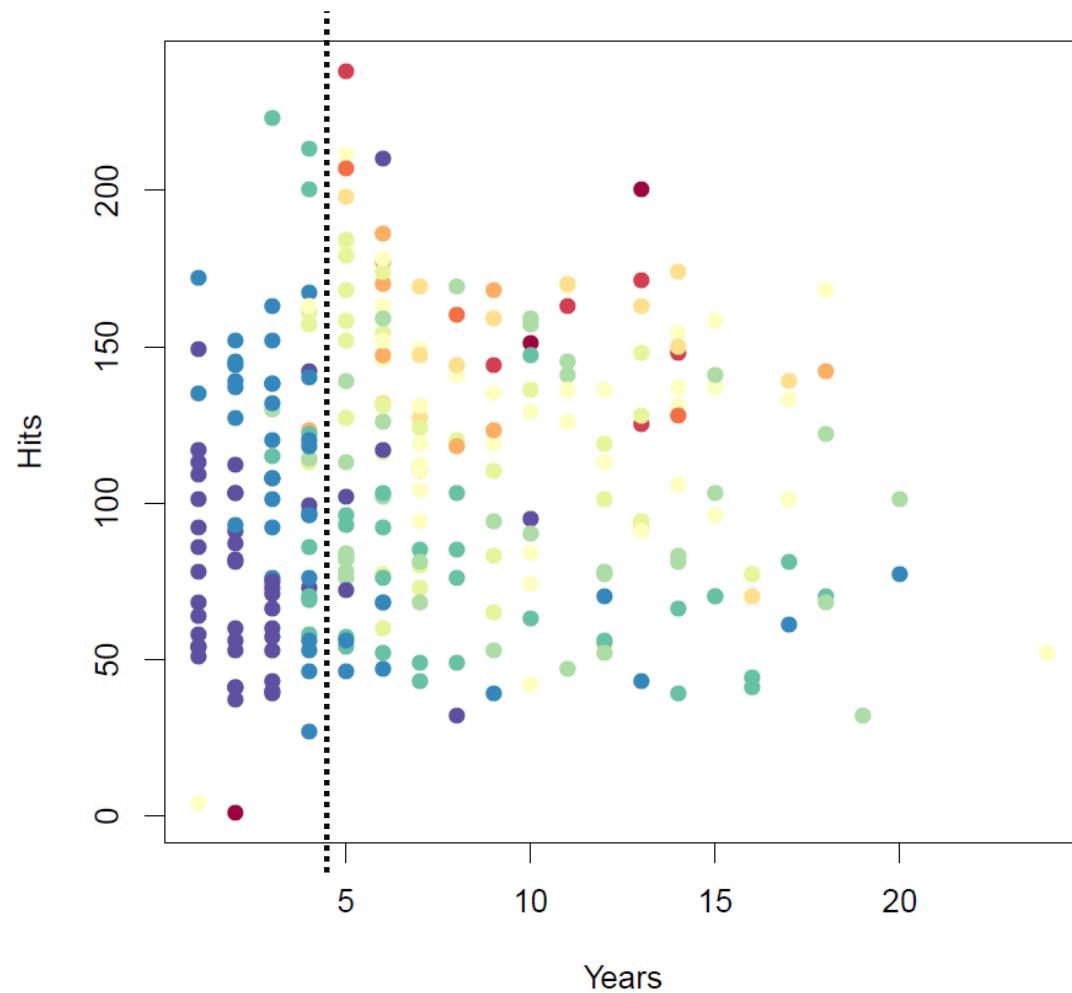
---

- Baseball salary data
  - Salary is color-coded from low (blue, green) to high (yellow, red)



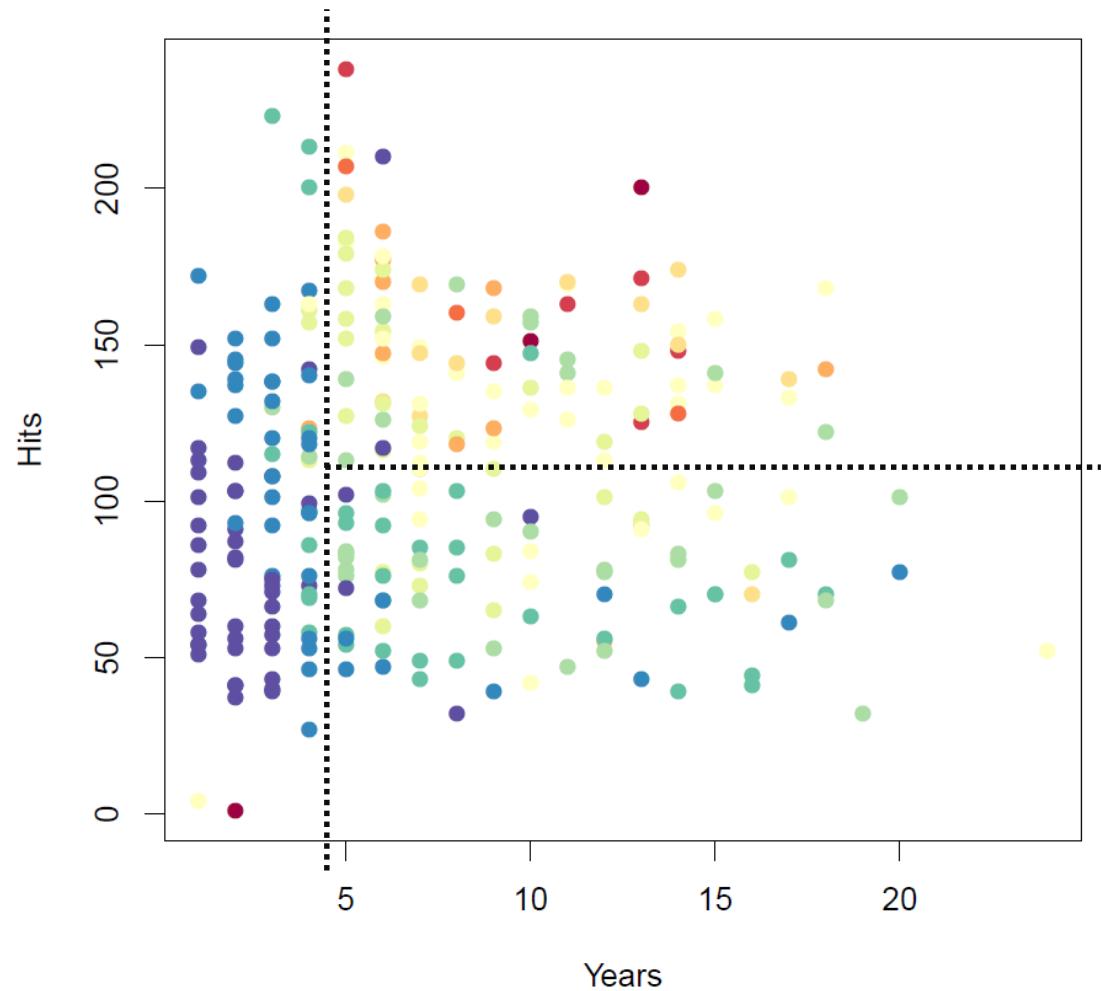
# Baseball salary data

---

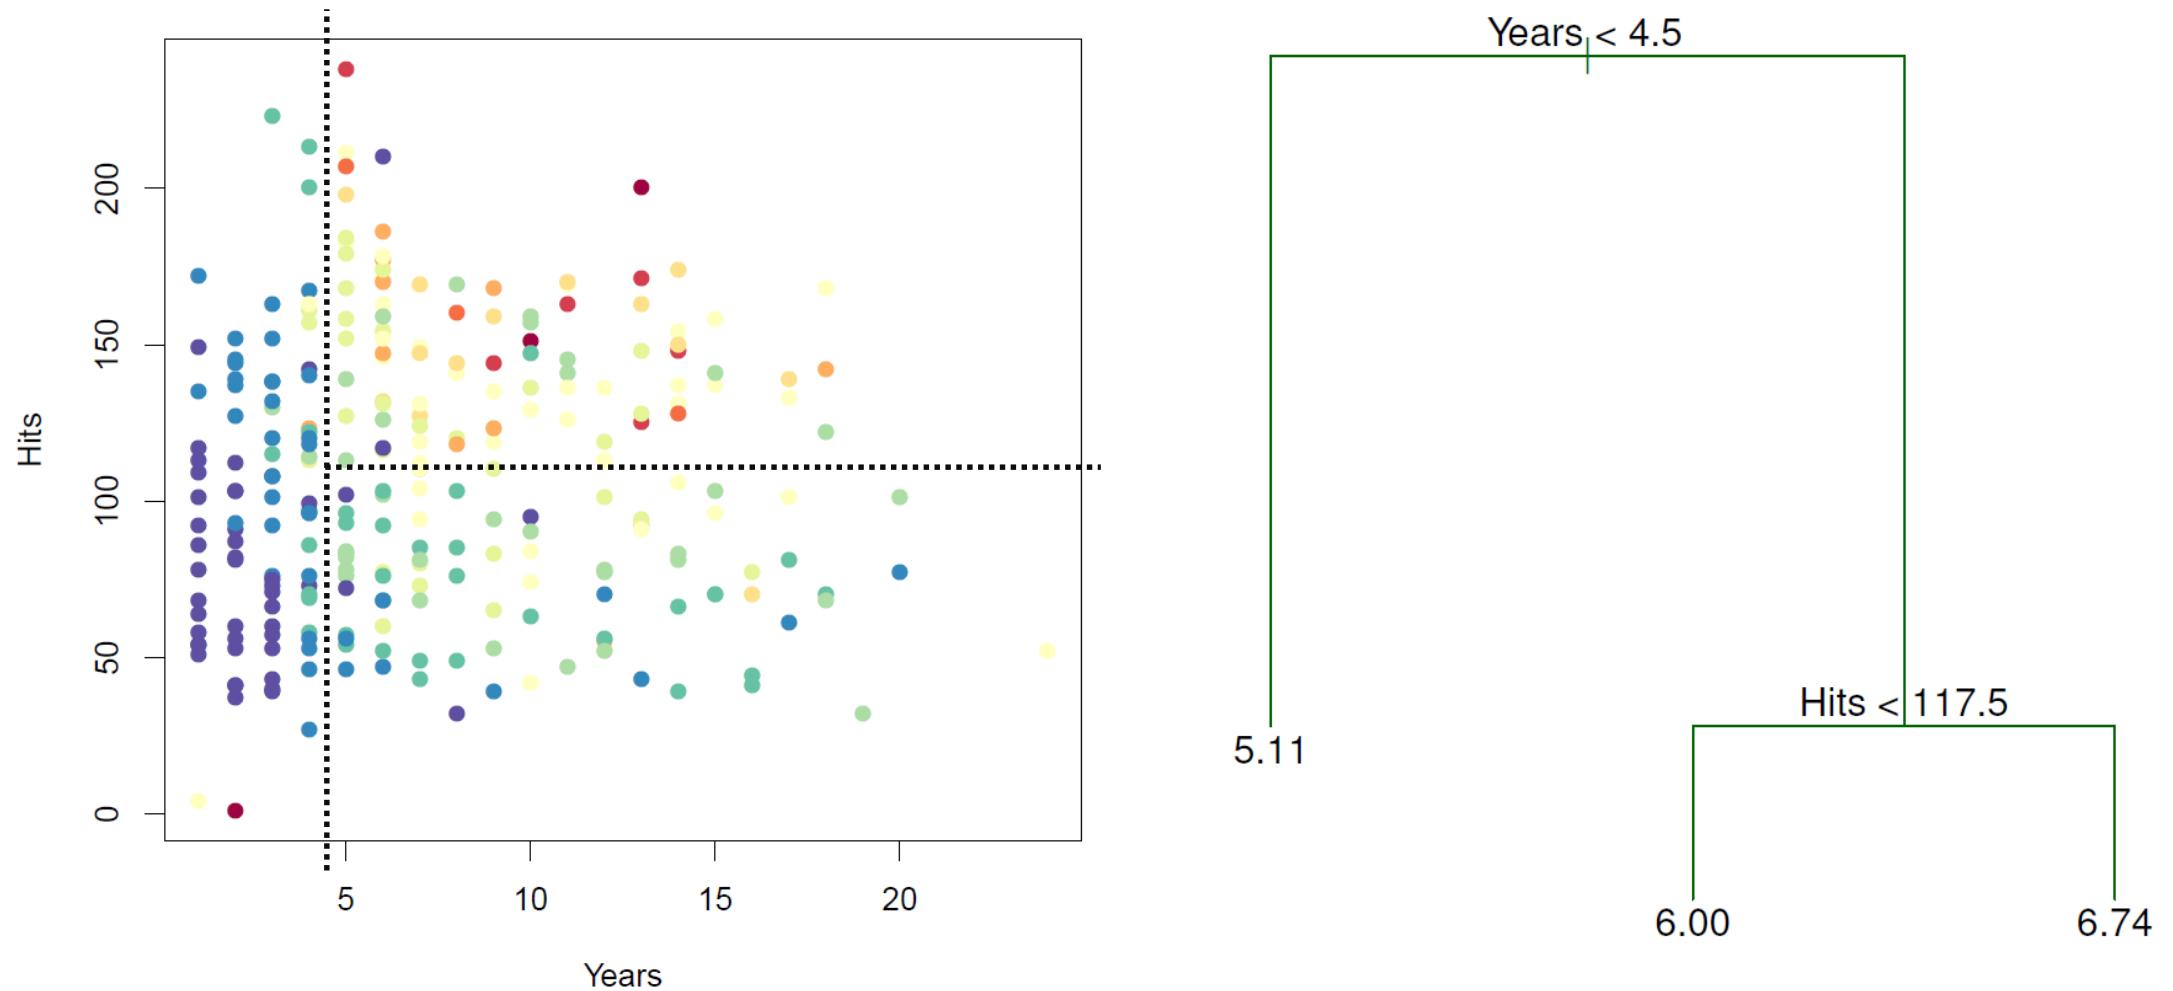


# Baseball salary data

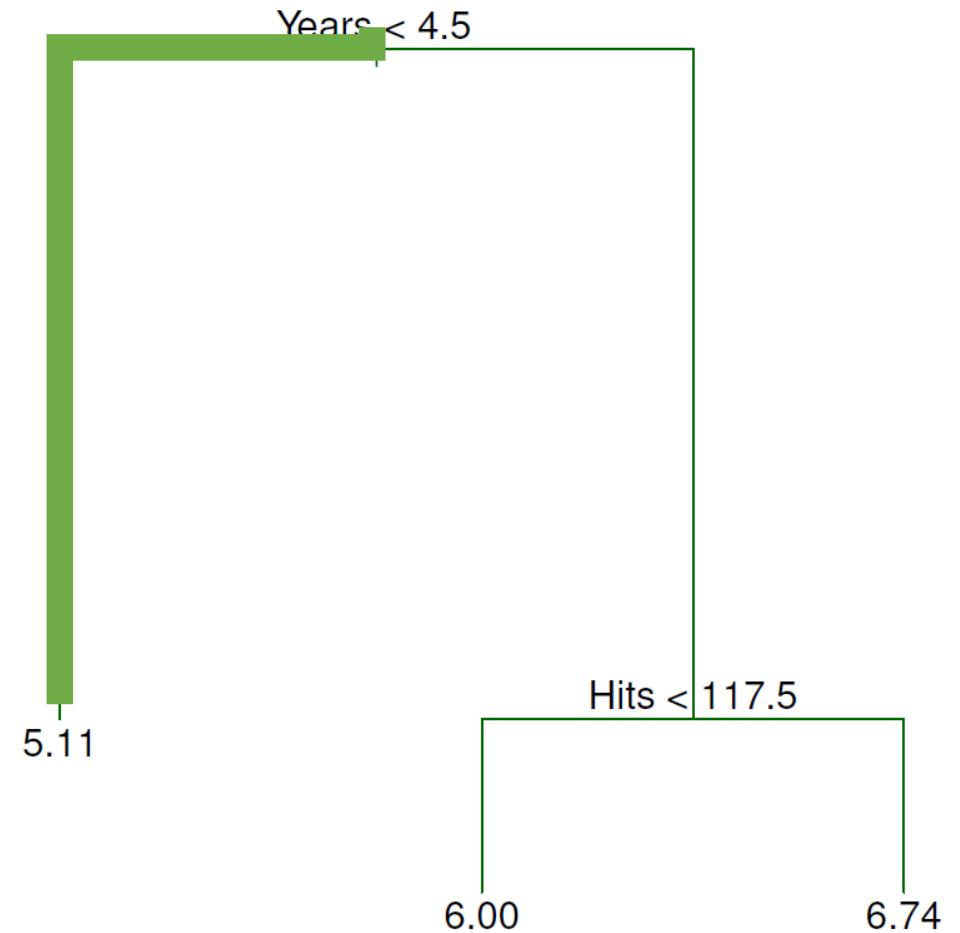
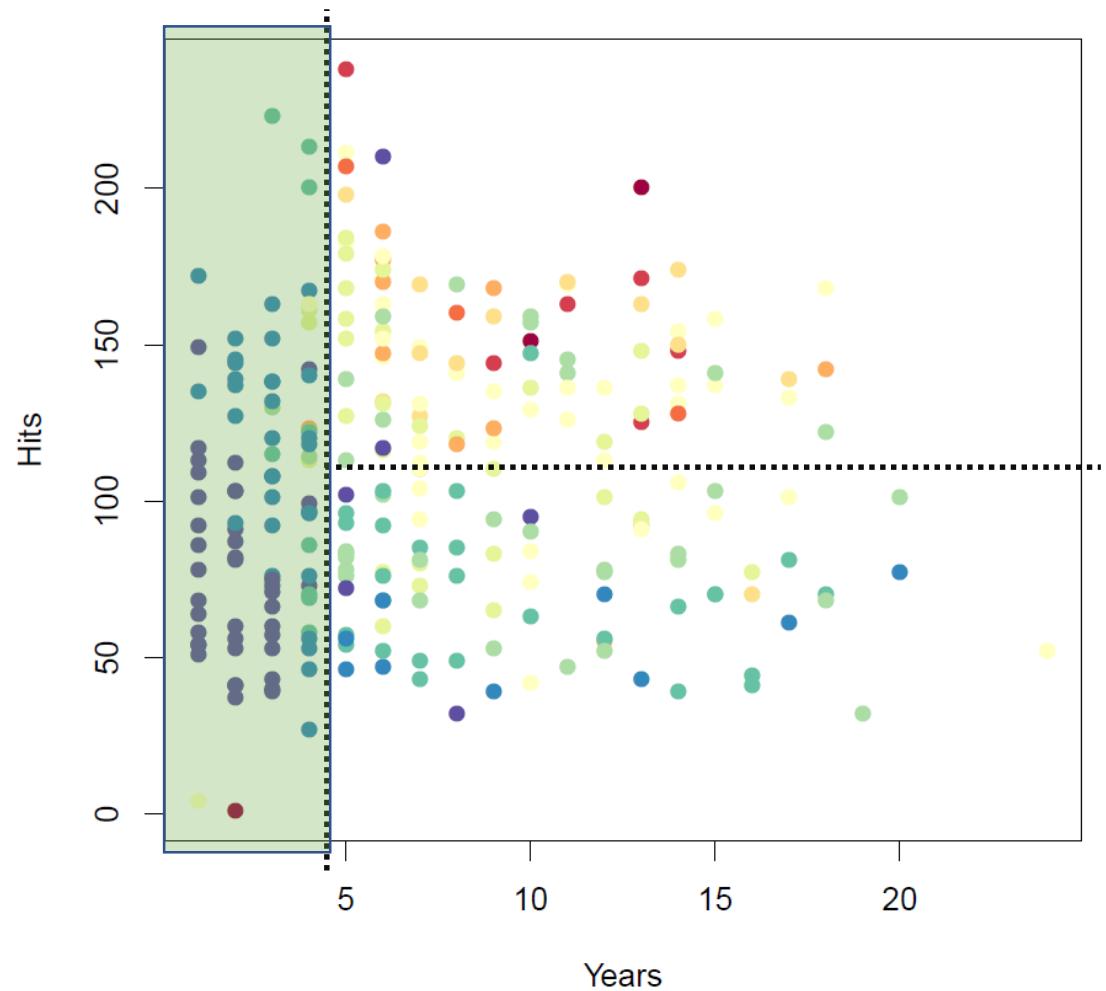
---



# Baseball salary data

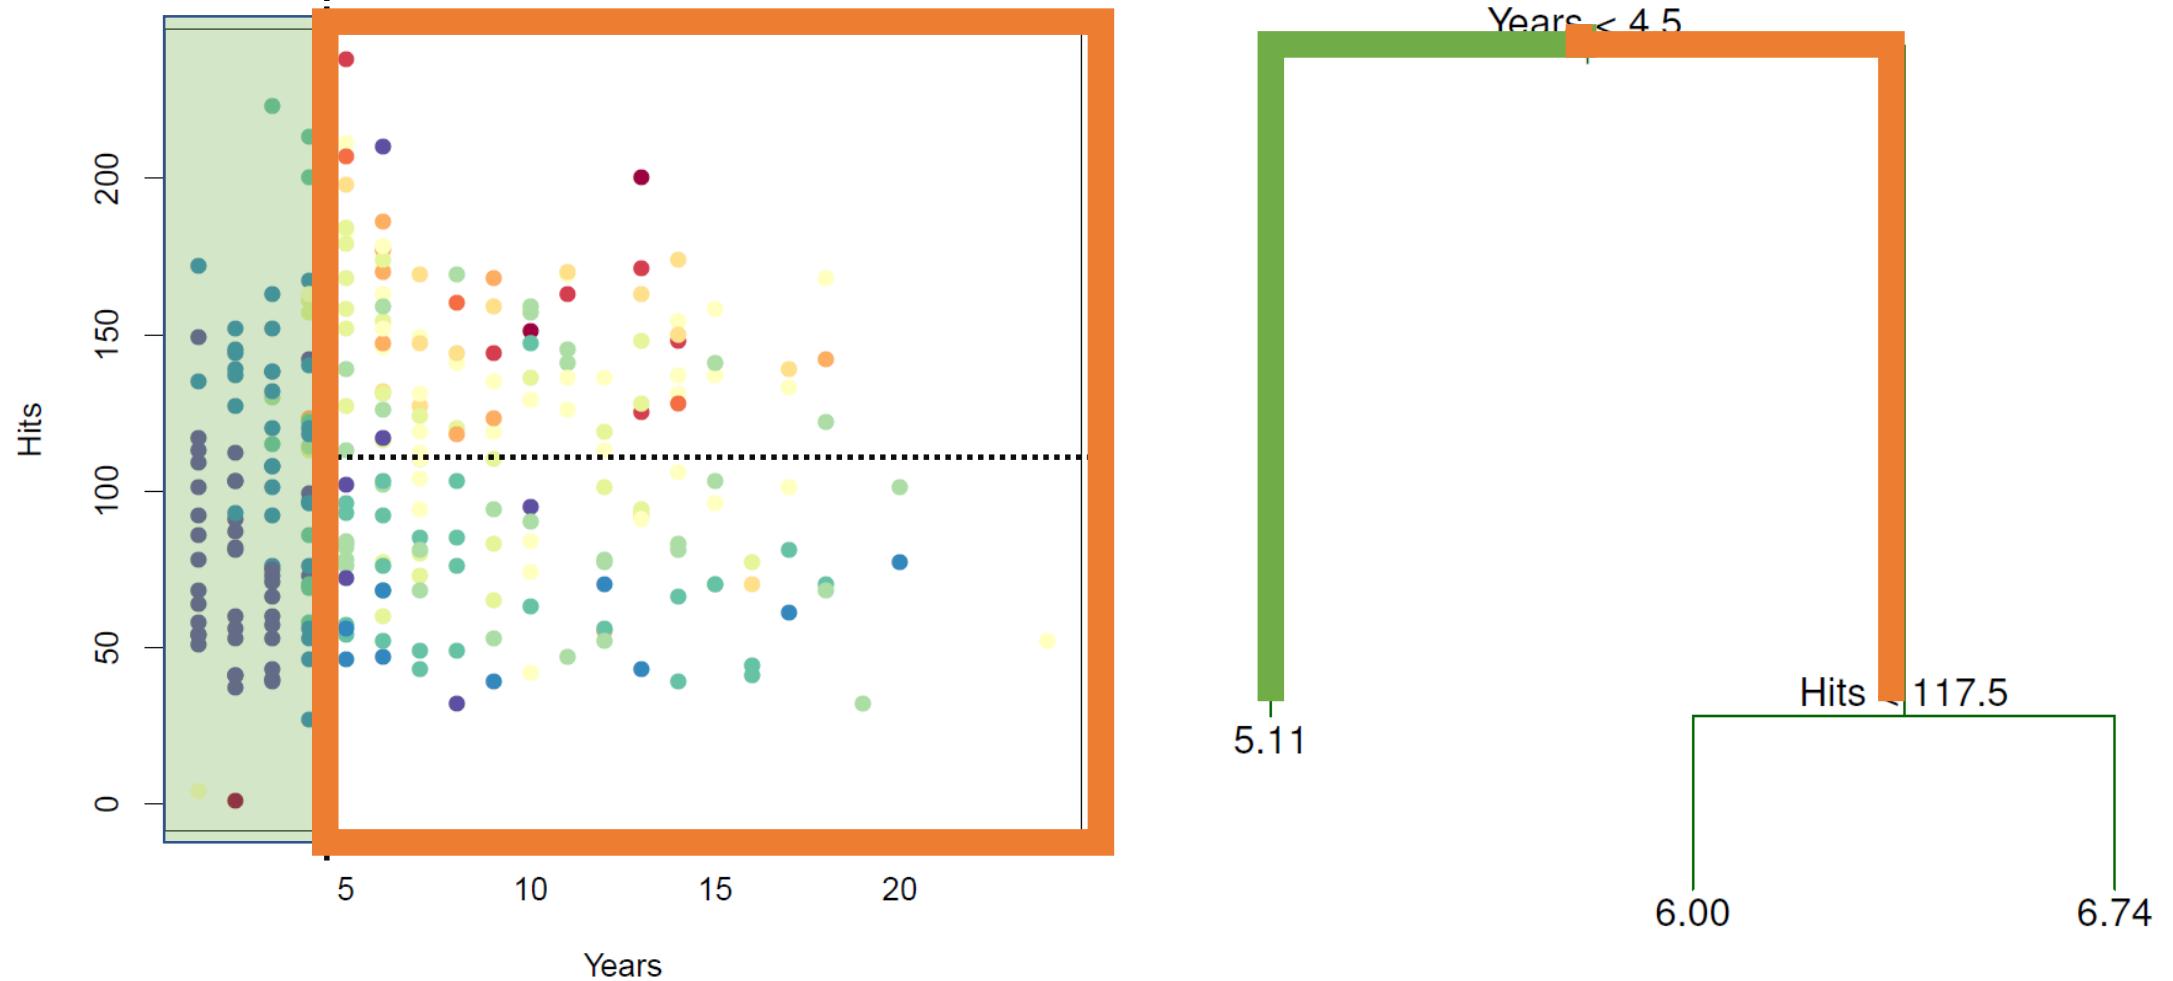


# Baseball salary data



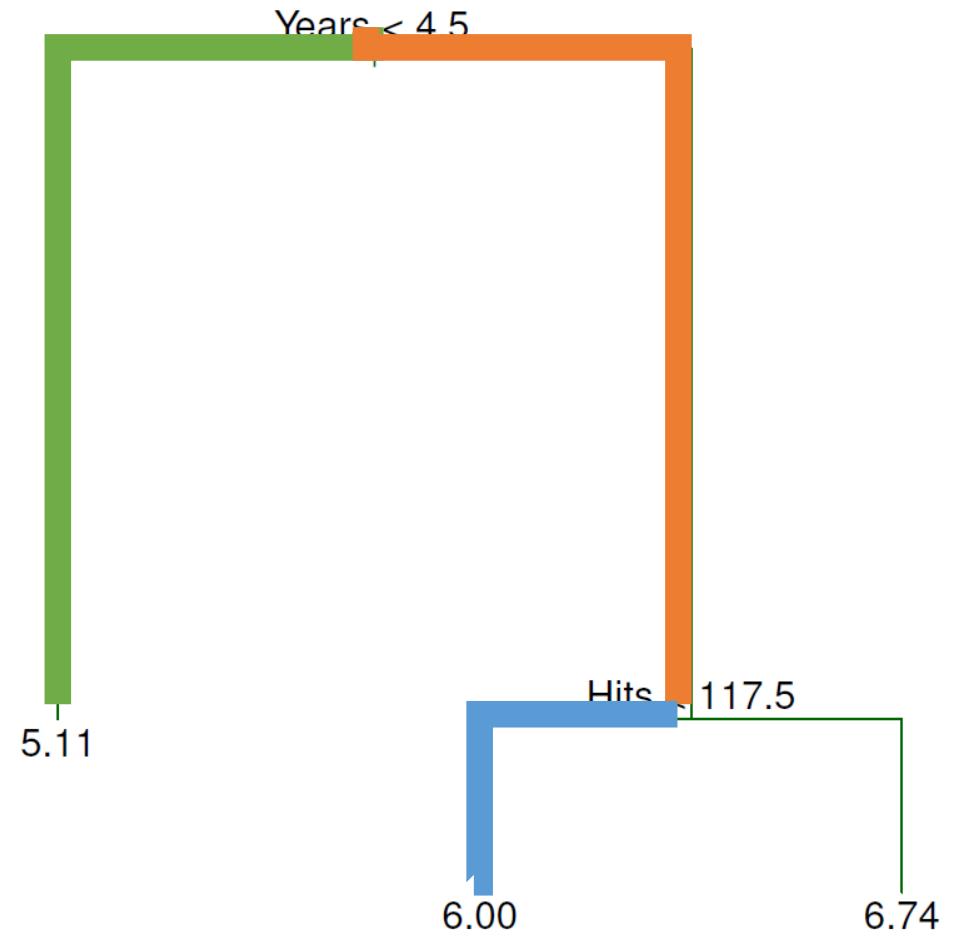
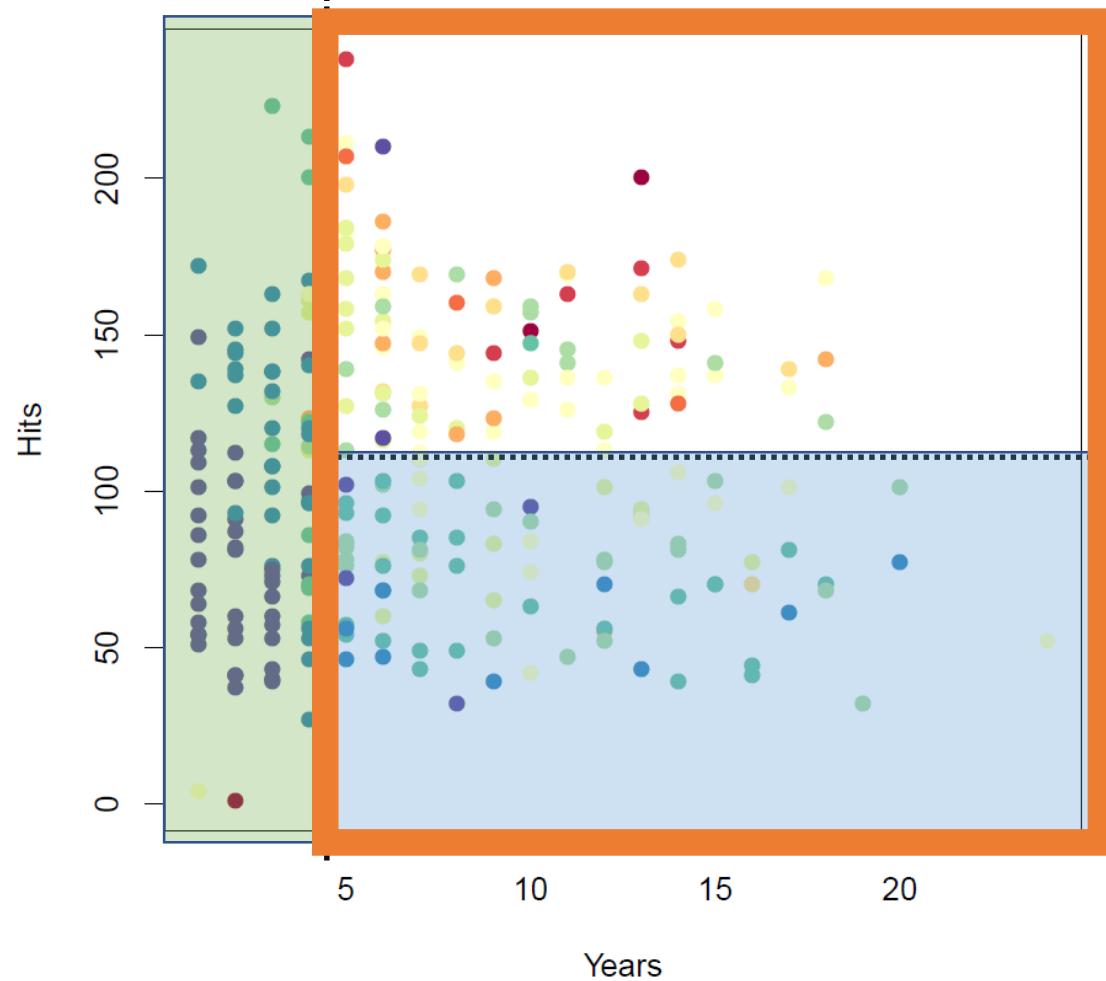
Hastie and Tibshirani

# Baseball salary data



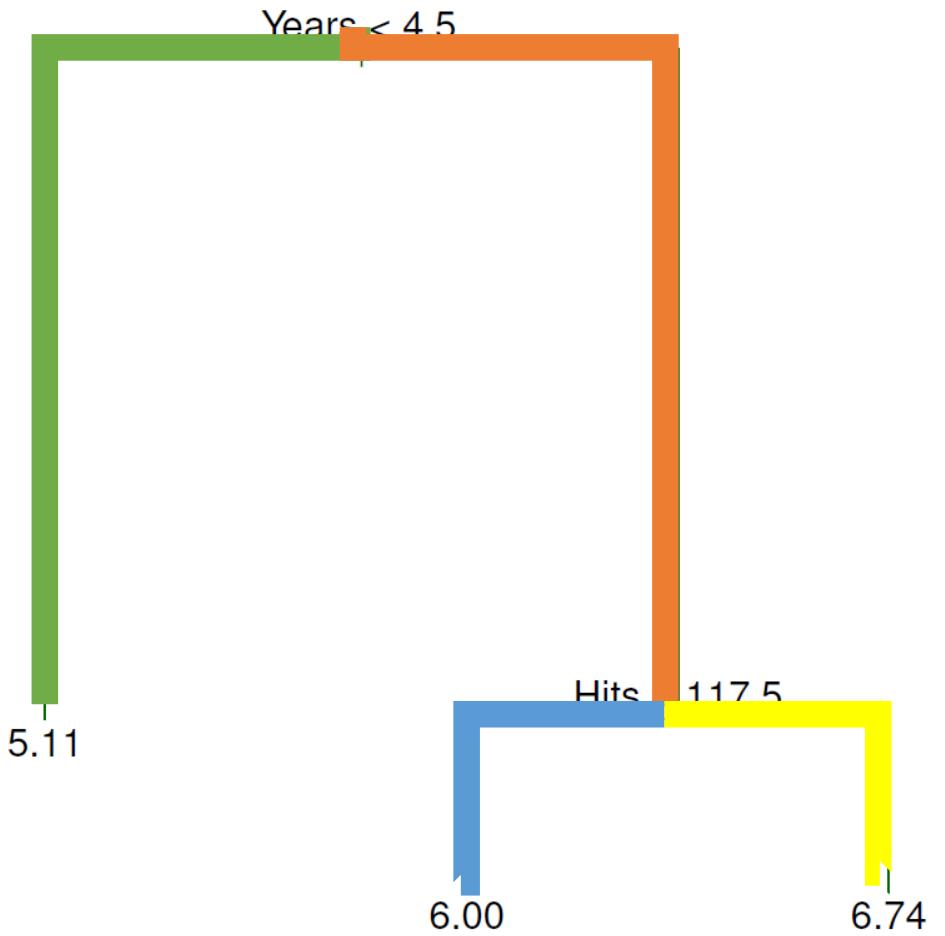
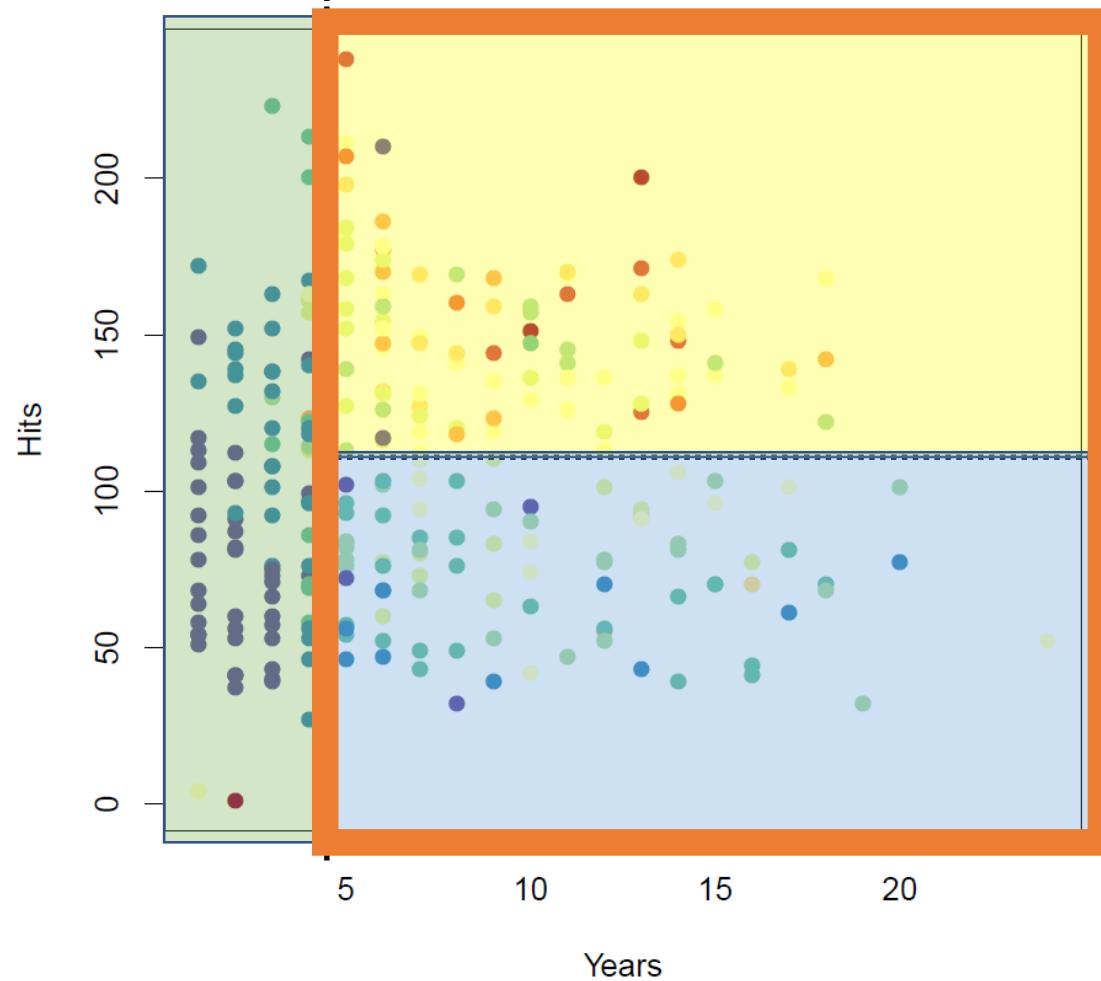
Hastie and Tibshirani

# Baseball salary data



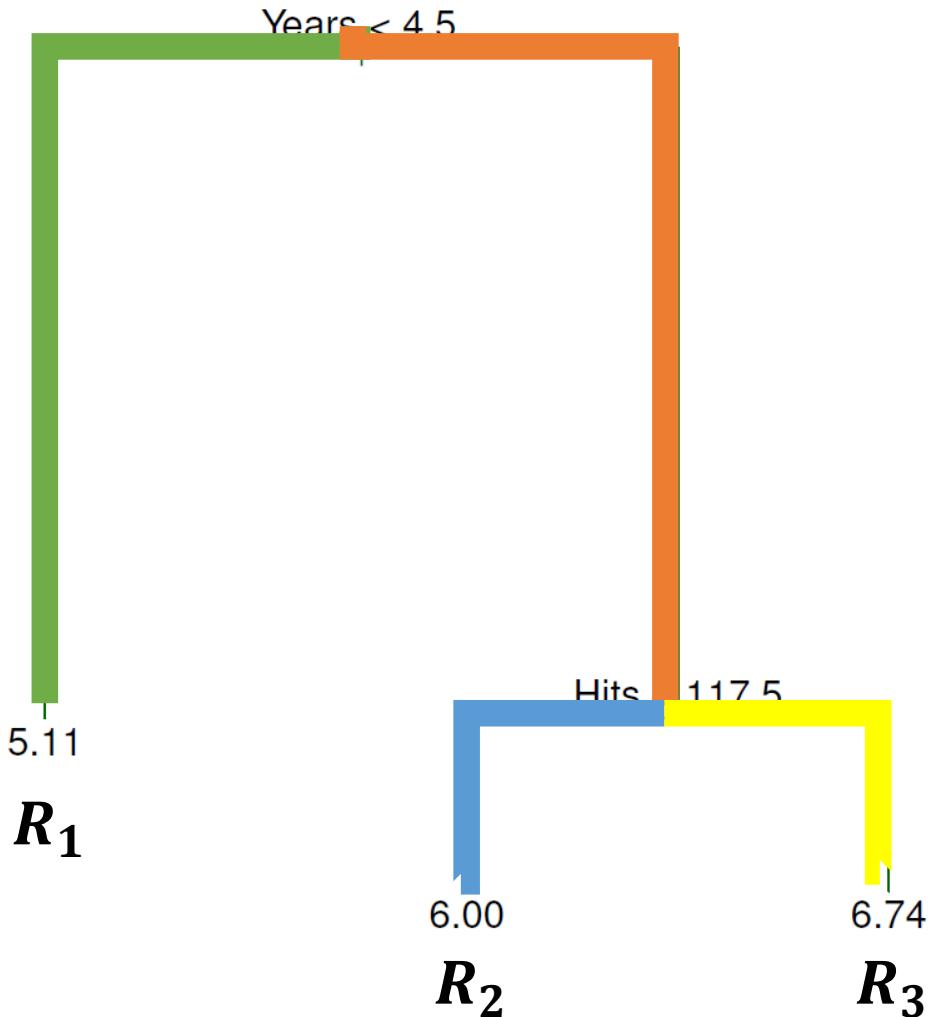
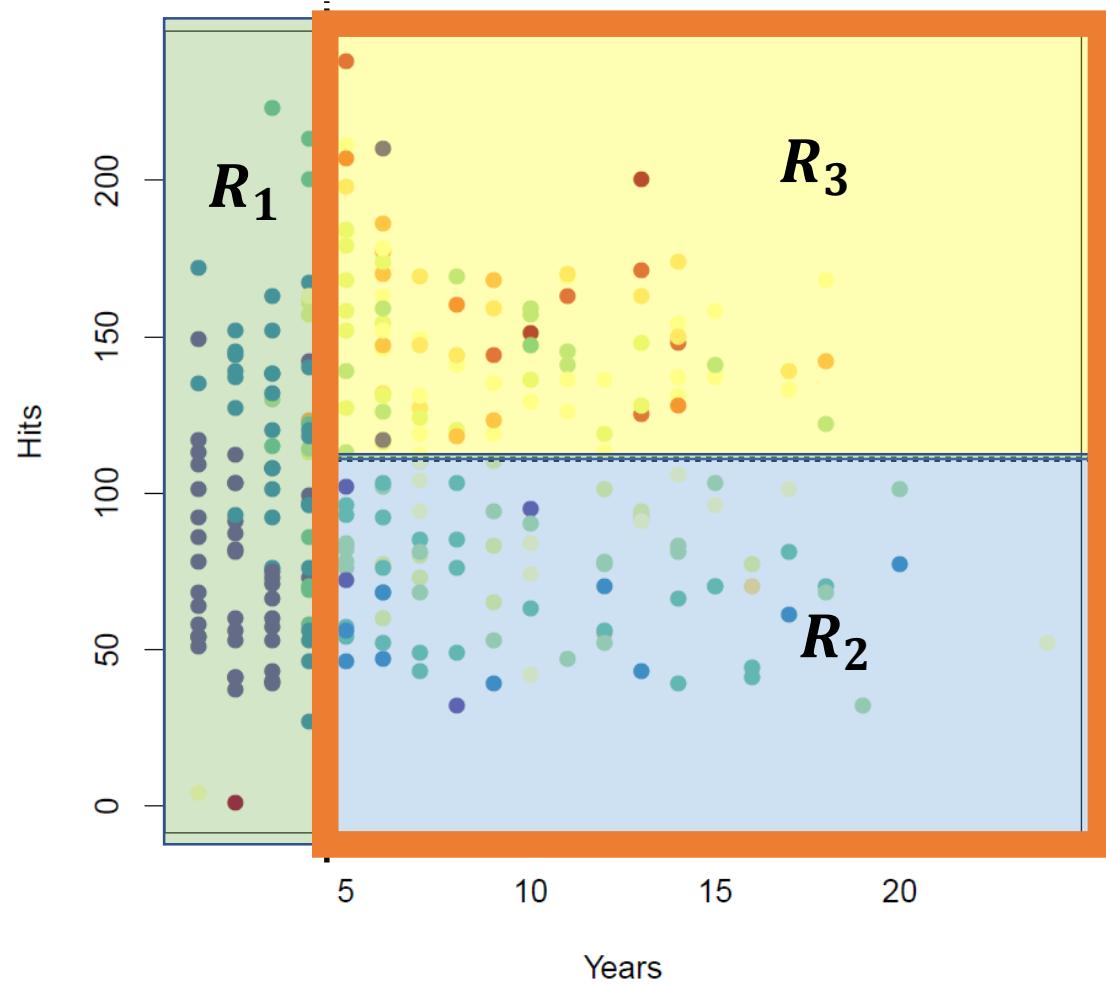
Hastie and Tibshirani

# Baseball salary data



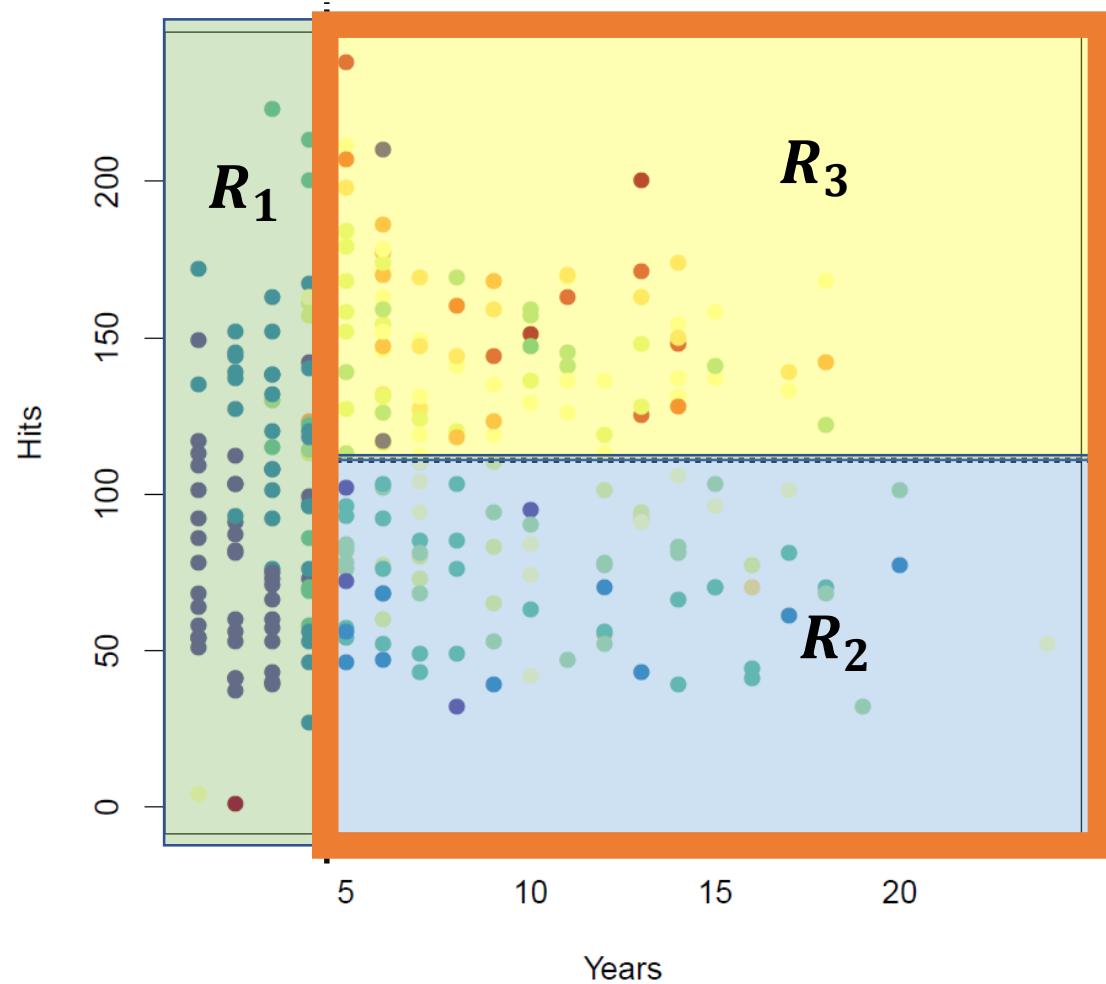
Hastie and Tibshirani

# Baseball salary data



# Baseball salary data

---



Overall the tree segments the feature space into three regions:

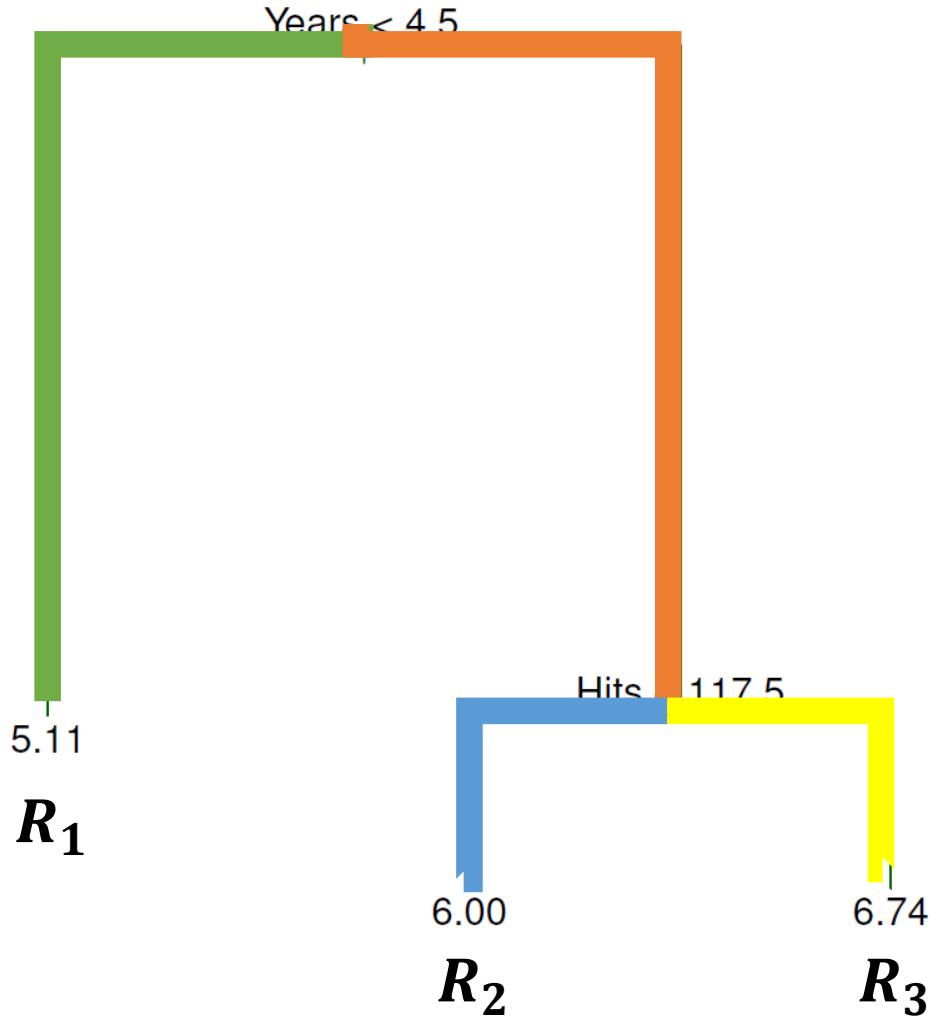
$$R_1 = \{x; \text{Years} < 4.5\}$$

$$R_2 = \{x; \text{Years} < 4.5, \text{Hits} < 117.5\}$$

$$R_3 = \{x; \text{Years} < 4.5, \text{Hits} \geq 117.5\}$$

# Baseball salary data

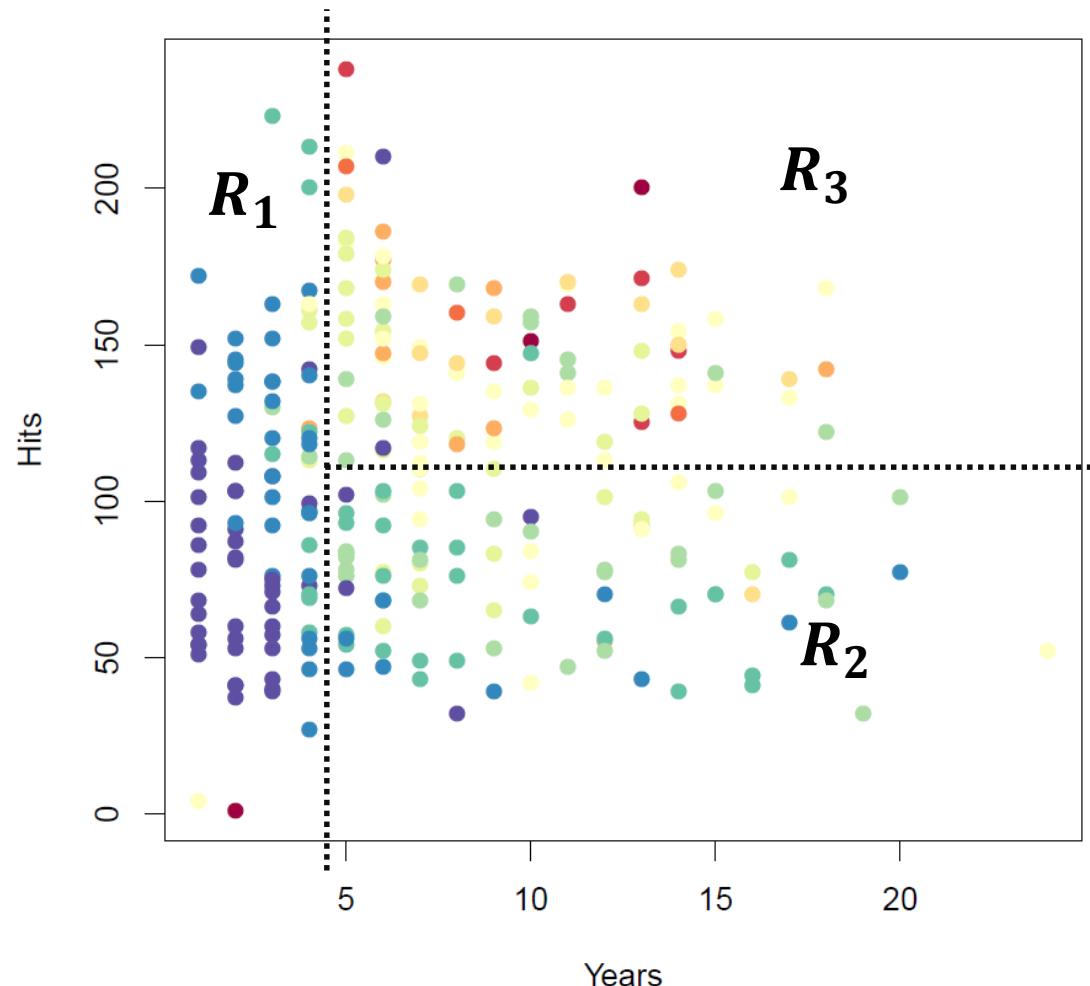
---



- Overall the tree segments the feature space into three regions:
- $R_1 = \{x; \text{Years} < 4.5\}$
- $R_2 = \{x; \text{Years} < 4.5, \text{Hits} < 117.5\}$
- $R_3 = \{x; \text{Years} < 4.5, \text{Hits} \geq 117.5\}$
- In tree terminology,  $R_1, R_2, R_3$  are known as **terminal nodes** whereas the nodes where the feature space is split are referred to as **internal nodes**.

# Interpretation of Results

---



- *Years* is the **most important factor** in determining salary. Players with less experience earn lower salaries than more experienced players.
- Given that a player is less experienced, the number of *Hits* that he made in the previous year seems to play little role in his salary.
- Among players who have been in the major leagues for five or more years, the number of *Hits* does affect salary

# Tree-building process - Regression

---

- Divide feature space into high-dimensional rectangles.
- The goal is to find rectangular regions  $R_1, R_2, \dots, R_J$  minimize error function:

$$E = \sum_{j=1}^J \sum_{i \in R_j} (x_i - \bar{x}_{R_j})^2$$

where  $\bar{x}_{R_j}$  is the mean of the training samples within region  $R_j$  .

# Tree-building process - Regression

---

- Divide feature space into high-dimensional rectangles.
- The goal is to find rectangular regions  $R_1, R_2, \dots, R_J$  minimize error function:

$$E = \sum_{j=1}^J \sum_{i \in R_j} (x_i - \bar{x}_{R_j})^2$$

where  $\bar{x}_{R_j}$  is the mean of the training samples within region  $R_j$ .

- Well... this is computationally infeasible! Thus we follow a greedy approach:
  - top-down approach: successively **split feature space**; each split is indicated via two new branches further down on the tree.
  - At each step, **best split is made at that particular step**, instead of looking ahead and picking a split that will lead to a better tree in some future step. (greedy approach)

# Tree-building process - 2

---

- First select the feature  $x_i$  and the cutpoint  $s$  in order to split the feature space into two regions:  $\{\mathbf{x}: x_i < s\}$  and  $\{\mathbf{x}: x_i \geq s\}$  for which  $E$  is minimized.

# Tree-building process - 2

---

- First select the feature  $x_i$  and the cutpoint  $s$  in order to split the feature space into two regions:  $\{\mathbf{x}: x_i < s\}$  and  $\{\mathbf{x}: x_i \geq s\}$  for which  $E$  is minimized.
- Repeat the process, i.e., best feature and best cutpoint to further minimize  $E$  (largest decrease) within each the segmented regions. Split the corresponding region.

# Tree-building process - 2

---

- First select the feature  $x_i$  and the cutpoint  $s$  in order to split the feature space into two regions:  $\{x: x_i < s\}$  and  $\{x: x_i \geq s\}$  for which  $E$  is minimized.
- Repeat the process, i.e., best feature and best cutpoint to further minimize  $E$  (largest decrease) within each the segmented regions. Split the corresponding region.
- Each time, we split one of the two previously identified regions.

# Tree-building process - 2

---

- First select the feature  $x_i$  and the cutpoint  $s$  in order to split the feature space into two regions:  $\{x: x_i < s\}$  and  $\{x: x_i \geq s\}$  for which  $E$  is minimized.
- Repeat the process, i.e., best feature and best cutpoint to further minimize  $E$  (largest decrease) within each the segmented regions. Split the corresponding region.
- Each time, we split one of the two previously identified regions.
- The whole process continues until a stopping criterion is reached.
  - A specific value for  $E$
  - number of samples in each regions above a threshold.

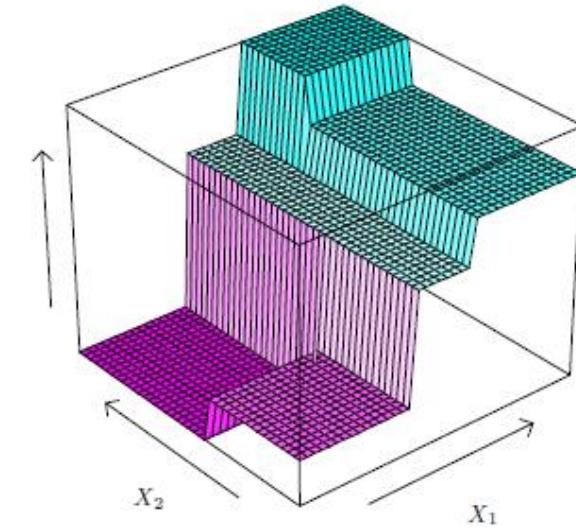
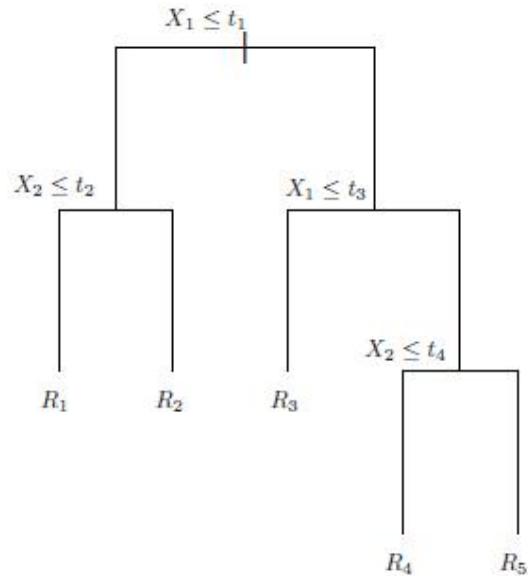
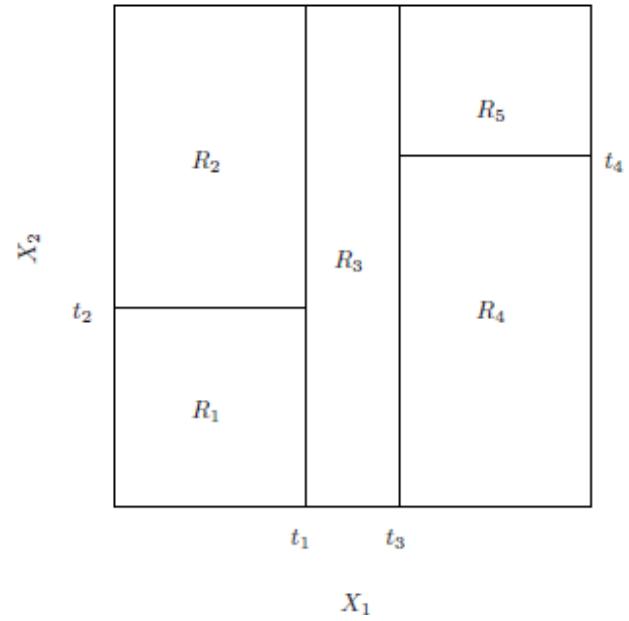
# Tree-building process - 2

---

- First select the feature  $x_i$  and the cutpoint  $s$  in order to split the feature space into two regions:  $\{\mathbf{x}: x_i < s\}$  and  $\{\mathbf{x}: x_i \geq s\}$  for which  $E$  is minimized.
  - Repeat the process, i.e., best feature and best cutpoint to further minimize  $E$  (largest decrease) within each the segmented regions. Split the corresponding region.
  - Each time, we split one of the two previously identified regions.
  - The whole process continues until a stopping criterion is reached.
    - A specific value for  $E$
    - number of samples in each regions above a threshold.
  - **Predict** the response for a given test sample by assigning the **mean** (regression) of the training samples in the region where the test sample belongs.
-

# Example

---



Hastie and Tibshirani

- **Left:** The output of recursive binary splitting on a two-dimensional example.
- **Center:** A tree corresponding to the partition in the top right panel.
- **Right:** A perspective plot of the prediction surface corresponding to that tree.

# Tree Pruning

---

- The abovementioned process may result in overfitting (really large trees). (remember with shallow trees I may get underfitted models).
- An alternate approach would be to grow a tree only so long as  $E$  decreases more than a certain threshold.
- This would result in medium sized trees but would cause the loss of larger decrease later on the training process.
  - seemingly worthless split early on in the tree might be followed by a very good split later on.
- A better strategy would be to grow a very large tree  $T_0$  and then prune it in order to obtain a **subtree that meet certain restrictions**.
  - Grow  $T_0$  and then prune it to  $T$  with a smaller number of terminal nodes  $|T|$ .

# Decision Tree – Classification ( $K$ classes)

---

- Same as in the regression case but we now split the feature space by minimizing

$$E = \sum_{j=1}^J Q_j$$

where  $Q_j$  can be one of the following options:

- Misclassification rate:  $Q_j = 1 - \max_k(p_{jk})$
- Negative cross entropy:  $Q_j = -\sum_{k=1}^K p_{jk} \ln p_{jk}$
- Gini index:  $Q_j = \sum_{k=1}^K p_{jk}(1 - p_{jk})$

where  $p_{jk}$  represents the proportion of training observations in the  $j$ th region that are from the  $k$ th class.

---

# Ensemble methods

---

- DT are easily interpretable and nice to visualize. It is said that DTs more closely **mirror human decision-making** than other ML approaches.
- Unfortunately, DT have **not the same level of predictive accuracy** as other classification approaches we have seen so far.
- Solution: Create ensembles of DTs!
  - Bagging/bootstrap aggregation with trees
  - Random Forests: bagging + random subspace methods
  - Boosting

# Random Forests

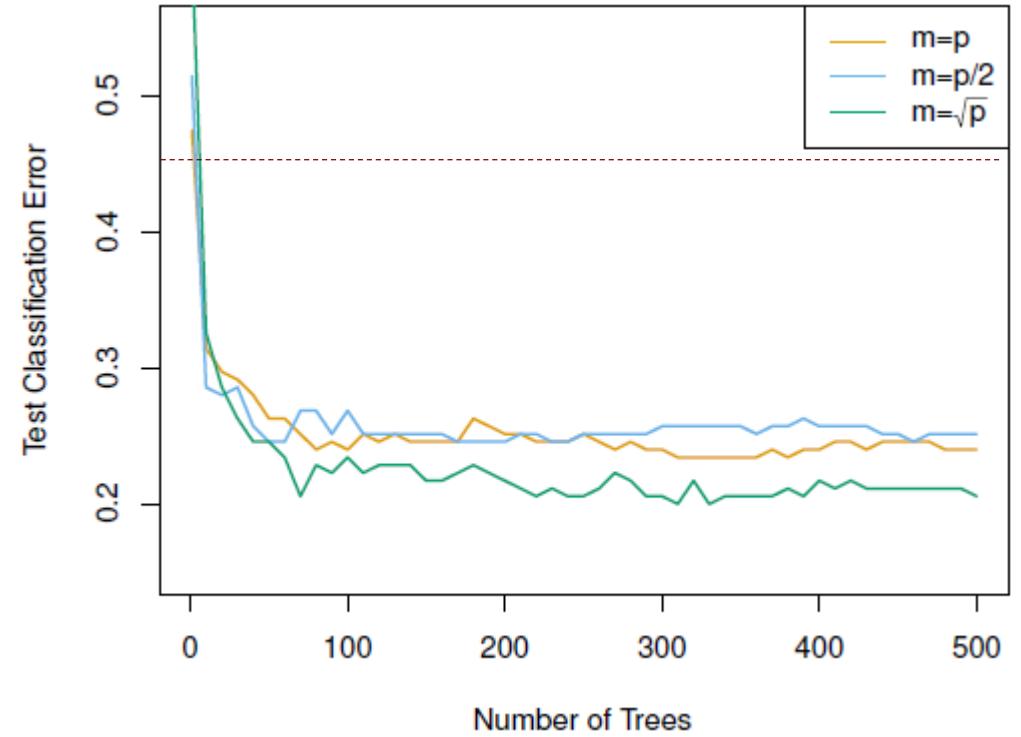
---

- Bootstrapping can lead to highly correlated DT.
  - If there are a few very strong features in the dataset all bootstrapped trees will use these features to split space.
- Solution:
  - Build an ensemble of trees by bootstrapping the dataset.
  - Use a feature bagging/random subspace method: for each tree, **every time a split is under consideration, a random selection of  $m$  out of  $p$  features is chosen** as a split candidate.
  - At each split a new selection is made

# Random Forests - Example

---

- Gene expression dataset
- Task: classify cancer type based on  $p = 500$  gene expressions (features).
- Random Forests ( $m < p$ )
- Just bootstrapping ( $m = p$ )
- Random Forests exhibit small improvement over just bootstrapping.
- Single tree performance (red dashed line) is much worse (45.7%)



Hastie and Tibshirani

# Sum up

---

- We can combine multiple models using different methods
  - Bootstrapping/Bagging
  - Random subspace methods/Feature Bagging
  - Boosting
- Trees are weak classifiers but with nice **interpretability**.
- Bagging, random forests and boosting are good methods for improving the prediction accuracy of trees.
- Random Forests and boosting are among the state-of-the-art methods for supervised learning but their results can be difficult to interpret.
  - The more complexity we add the more difficult it becomes to interpret our results.



ARISTOTLE UNIVERSITY OF THESSALONIKI



FACULTY OF ENGINEERING

# Questions?

*Pattern Recognition & Machine Learning*  
*Model Combination - Trees*